

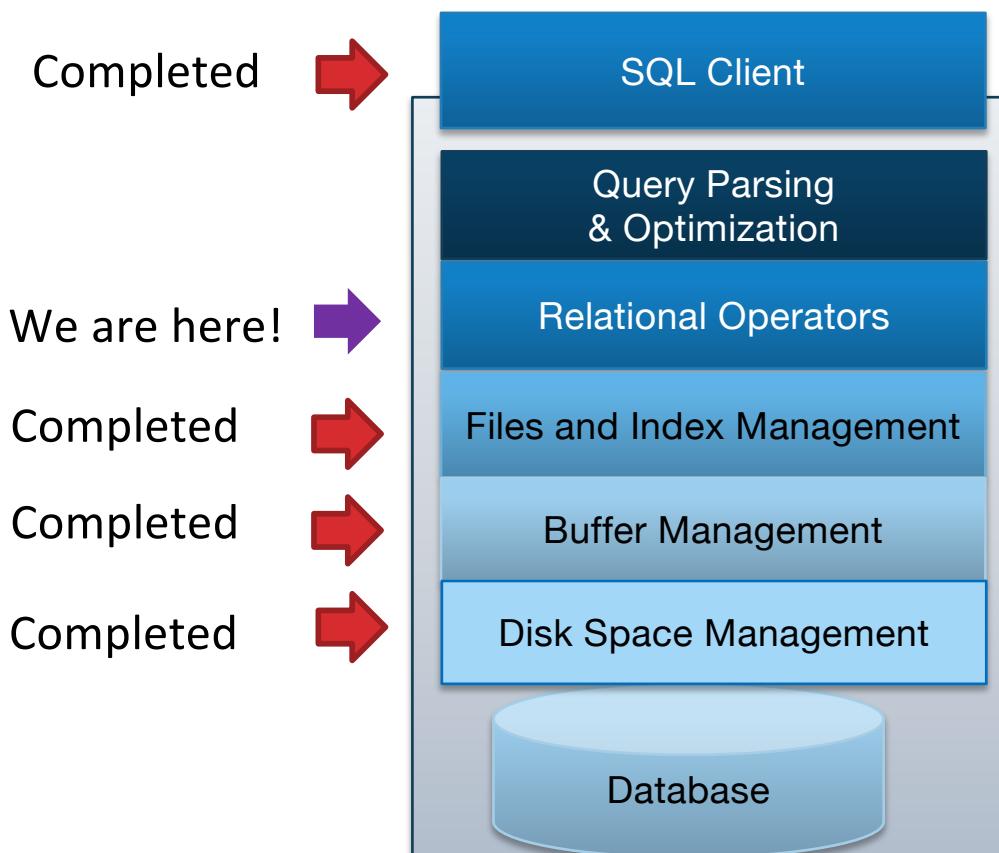
Relational Algebra

Alvin Cheung
Fall 2022

Reading: R&G, Chapters 4.1 - 4.2



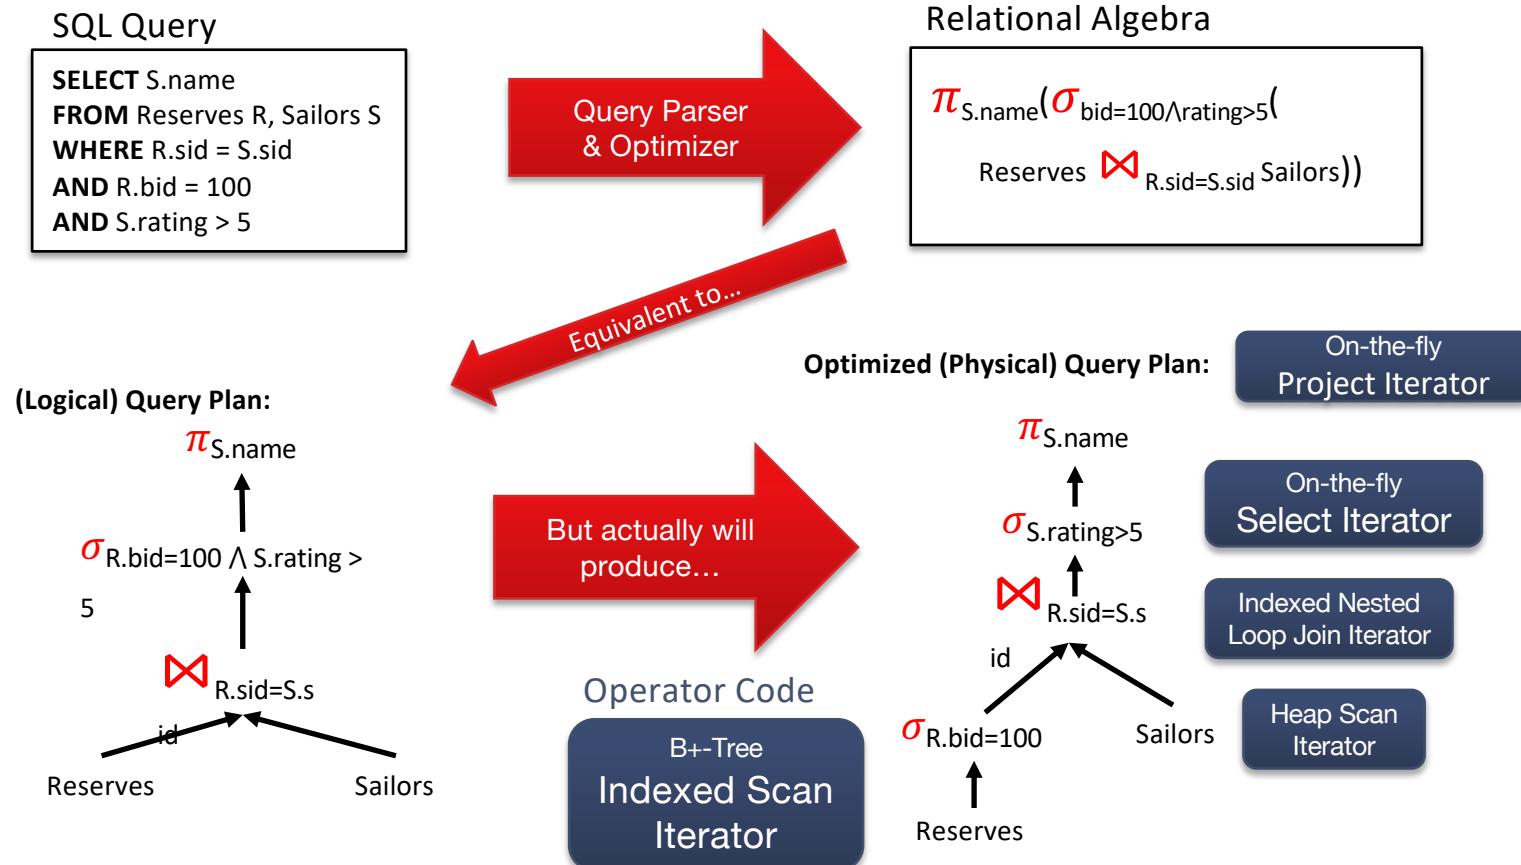
Architecture of a DBMS: What we've learned



Today: *definitions* of the relational operators.

Coming soon:
optimization

An Overview of the Layer Above



SQL vs Relational Algebra



SQL Query

```
SELECT S.name  
FROM Reserves R, Sailors S  
WHERE R.sid = S.sid  
AND R.bid = 100  
AND S.rating > 5
```

Query Parser
& Optimizer

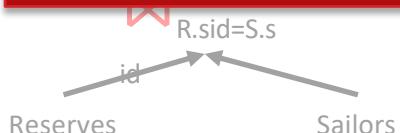
Relational Algebra

$$\pi_{S.name}(\sigma_{bid=100 \wedge rating>5}(Reserves \bowtie_{R.sid=S.sid} Sailors))$$

Equivalent to...

SQL

A **declarative** expression
of the query result



But actually will
produce...

Relational Algebra

Operational description of
a computation.

Operator Code
B+-Tree
Indexed
Iterator

On-the-fly
Optimizer

Heap Scan

Systems execute relational algebra
query plan.

SQL vs. Relational Algebra



SELECT S.name

FROM Reserves R, Sailors S

WHERE R.sid = S.sid

AND R.bid = 100

AND S.rating > 5

$\pi_{S.name}(\sigma_{R.bid=100 \wedge S.rating>5}(R \bowtie_{R.sid=S.sid} S))$

- *Why humans like SQL*
 - It's declarative
 - Say **what** you want, not **how** to get it
 - Enables system to optimize the **how**
- *Why systems like relational algebra*
 - It's operational
 - It describes the steps for **how** to compute a query result
- DBMSs transform SQL into relational algebra expressions, and find the best implementation to compute query result

Relational Algebra Preliminaries



- Algebra of operators on relation instances
 - Just like other algebras: linear algebra or elementary algebra
 - Operating on matrices or variables
- $\pi_{S.name}(\sigma_{R.bid=100 \wedge S.rating>5}(R \bowtie_{R.sid=S.sid} S))$
- Closed: result is also a relation instance
 - Enables rich composition!
 - Just like a linear algebraic expression on matrices returns a matrix
- Typed: input schema determines output
 - Can statically check whether queries are legal.
 - Same story for linear algebra – input sizes determine output sizes

Relational Algebra and Sets



- Pure relational algebra has set semantics
 - No duplicate tuples in a relation instance
 - vs. SQL, which has multiset (bag) semantics
 - We will switch to multiset in the system discussion

Relational Algebra Operators: Unary



- Unary Operators: on **single relation**
- **Projection** (π): Retains only desired columns (vertical)
- **Selection** (σ): Selects a subset of rows (horizontal)
- **Renaming** (ρ): Rename attributes and relations.

Relational Algebra Operators: Binary



- Binary Operators: on **pairs of relations**
- **Union** (\cup): Tuples in r_1 or in r_2 .
- **Set-difference** ($-$): Tuples in r_1 , but not in r_2 .
- **Cross-product** (\times): Allows us to combine two relations.

Relational Algebra Operators: Compound



- Compound Operators: ops that can be expressed using the 6 basic ops earlier
- **Intersection** (\cap): Tuples in r1 and in r2.
- **Joins** (\bowtie_θ , \bowtie): Combine relations that satisfy predicates

Relational Algebra Operators



- **Projection** (π): Retains only desired columns (vertical)
- **Selection** (σ): Selects a subset of rows (horizontal)
- **Renaming** (ρ): Rename attributes and relations
- **Union** (\cup): Tuples in r1 or in r2.
- **Set-difference** ($-$): Tuples in r1, but not in r2.
- **Cross-product** (\times): Allows us to combine two relations.
- **Intersection** (\cap): Tuples in r1 and in r2.
- **Joins** (\bowtie_θ , \bowtie): Combine relations that satisfy predicates

Projection (π)

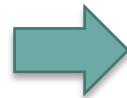
- Corresponds to the SELECT list in SQL
- Schema determined by schema of attribute list
 - Names and types correspond to input attributes
- Selects a subset of columns (vertical)

$$\pi_{\text{sname}, \text{age}}(\text{S2})$$

table
 —————
 Relational Instance S2

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 28 | yuppy | 9 | 35.0 |
| 31 | lubber | 8 | 55.5 |
| 44 | guppy | 5 | 35.0 |
| 58 | rusty | 10 | 35.0 |

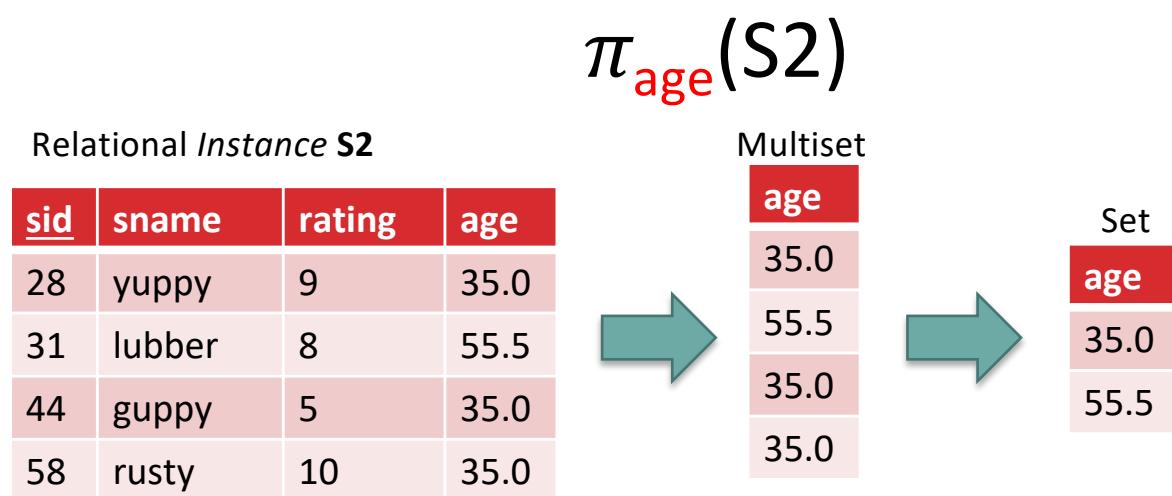
List of Attributes



| sname | age |
|--------|------|
| yuppy | 35.0 |
| lubber | 55.5 |
| guppy | 35.0 |
| rusty | 35.0 |

Projection (π), cont.

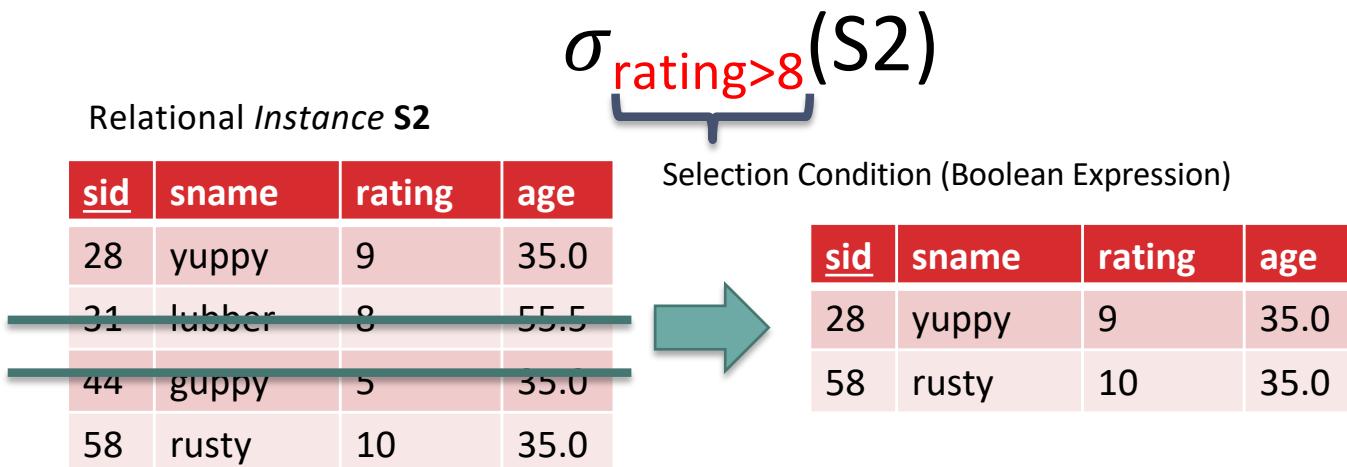
- Set semantics results in fewer rows
 - Real systems don't automatically remove duplicates
 - Why? (Semantics and Performance reasons)



Selection(σ)



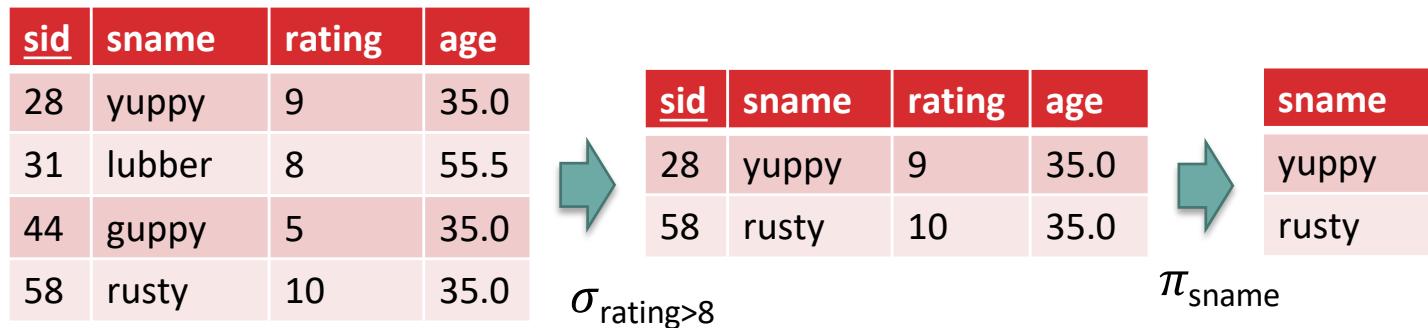
- Selects a subset of rows (horizontal)
- Corresponds to the WHERE clause in SQL
- Output schema same as input
- Duplicate Elimination? Not needed if input is a set.



Composing Select and Project



- Names of sailors with rating > 8: $\pi_{\text{pname}}(\sigma_{\text{rating}>8}(\text{S2}))$



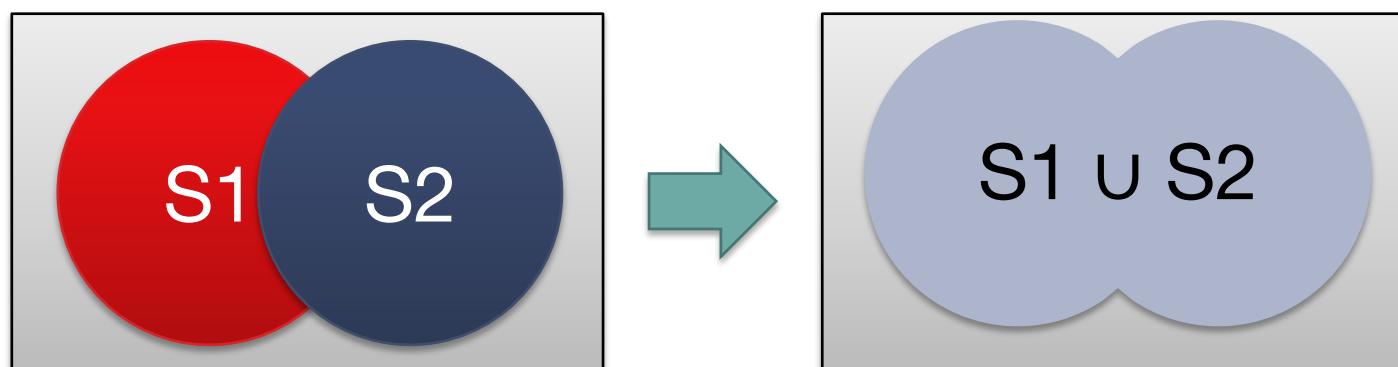
- What about: $\sigma_{\text{rating}>8}(\pi_{\text{sname}}(\text{S2}))$
 - Invalid types. Input to $\sigma_{\text{rating}>8}$ does not contain rating.

Union (\cup)



- *Takes the set union of two sets*
- The two input relations must be *compatible*:
 - Same sequence of attributes and types thereof
- SQL Expression: UNION

$S1 \cup S2$



Union (\cup) VS Union ALL



- In union under set semantics, duplicate elimination is needed
- SQL Expression: UNION (get rid of duplicates) vs. UNION ALL (keep dup.)

Relational Instance S1

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

Relational Instance S2

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 28 | yuppy | 9 | 35.0 |
| 31 | lubber | 8 | 55.5 |
| 44 | guppy | 5 | 35.0 |
| 58 | rusty | 10 | 35.0 |

$S1 \cup S2$

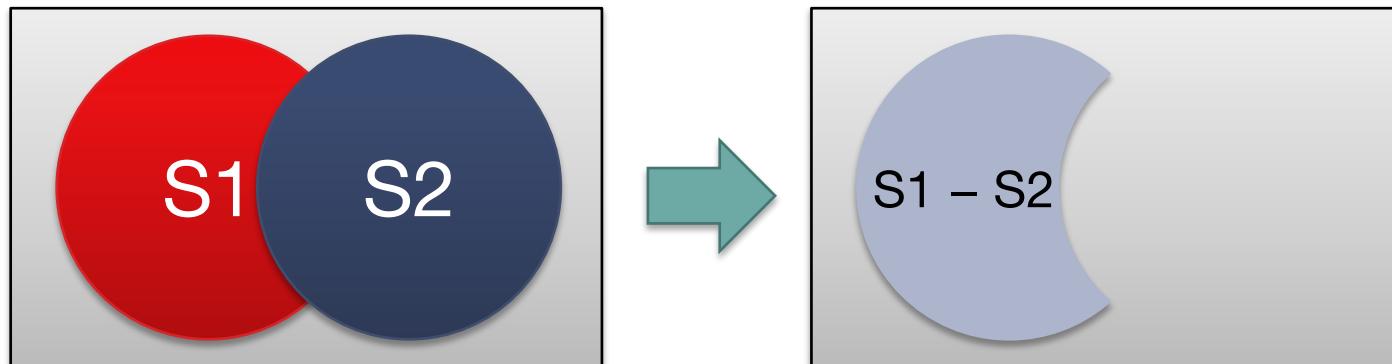
| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 22 | dustin | 7 | 45 |
| 28 | yuppy | 9 | 35.0 |
| 31 | lubber | 8 | 55.5 |
| 44 | guppy | 5 | 35.0 |
| 58 | rusty | 10 | 35.0 |

Set Difference (-)



- Same as with union, both input relations must be *compatible*.
- SQL Expression: EXCEPT *(distinct rows of left table)*

$S1 - S2$



Set Difference (-), cont.



- Q: Do we need to eliminate duplicates like in UNION?
 - Not required if inputs are sets
- SQL Expression: EXCEPT vs EXCEPT ALL**
 - Same as UNION/UNION ALL
 - In EXCEPT duplicates are eliminated if they exist

(allows duplicate
from left table)

S1 – S2

Relational Instance S1

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

Relational Instance S2

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 28 | yuppy | 9 | 35.0 |
| 31 | lubber | 8 | 55.5 |
| 44 | guppy | 5 | 35.0 |
| 58 | rusty | 10 | 35.0 |

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|-----|
| 22 | dustin | 7 | 45 |

S2 – S1

| <u>sid</u> | sname | rating | age |
|------------|-------|--------|------|
| 28 | yuppy | 9 | 35.0 |
| 44 | guppy | 5 | 35.0 |

Cross-Product (\times)



- **R1 \times S1: Each row of R1 paired with each row of S1**

R1:

| sid | bid | day |
|-----|-----|----------|
| 22 | 101 | 10/10/96 |
| 58 | 103 | 11/12/96 |

S1:

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

\times

=

| sid | bid | day | sid | sname | rating | age |
|-----|-----|----------|-----|--------|--------|------|
| 22 | 101 | 10/10/96 | 22 | dustin | 7 | 45.0 |
| 22 | 101 | 10/10/96 | 31 | lubber | 8 | 55.5 |
| 22 | 101 | 10/10/96 | 58 | rusty | 10 | 35.0 |
| 58 | 103 | 11/12/96 | 22 | dustin | 7 | 45.0 |
| 58 | 103 | 11/12/96 | 31 | lubber | 8 | 55.5 |
| 58 | 103 | 11/12/96 | 58 | rusty | 10 | 35.0 |

- How many rows in result?
 - $|R1|^*|S1|$
- Do we need to worry about schema compatibility?
 - Not needed.
- Do we need to do duplicate elimination?
 - None generated.

renaming columns

Renaming (ρ = “rho”)



- Renames relations and their attributes
- Convenient to avoid confusion when two relations overlap in attributes
- Can omit output name if we don't want to rename the output

$$\rho_{R(\underline{sid2}, \underline{bid2}, \underline{day})}(R1)$$

or

$$\rho_{R(sid \rightarrow sid2, bid \rightarrow bid2)}(R1)$$

R1:

| <u>sid</u> | <u>bid</u> | <u>day</u> |
|------------|------------|------------|
| 22 | 101 | 10/10/96 |
| 58 | 103 | 11/12/96 |

R:

| <u>sid2</u> | <u>bid2</u> | <u>day</u> |
|-------------|-------------|------------|
| 22 | 101 | 10/10/96 |
| 58 | 103 | 11/12/96 |

Naming tables/relations

Renaming (ρ = “rho”) contd.



- Yet another shorthand for renaming

$$\rho_{\text{Temp1}(\text{R1.sid} \rightarrow \text{sid1}, \text{S1.sid} \rightarrow \text{sid2})} (\text{R1} \times \text{S1})$$

↓ ↓ ↓
 Output Name List of cols to rename Input Relation

- Again, can omit output name if we don't want to rename the output
- For this case, we can equivalently rename each relation and then do cross-product

$\text{R1} \times \text{S1}$

| sid | bid | day | sid | sname | rating | age |
|-----|-----|----------|-----|--------|--------|------|
| 22 | 101 | 10/10/96 | 22 | dustin | 7 | 45.0 |
| 22 | 101 | 10/10/96 | 31 | lubber | 8 | 55.5 |
| 22 | 101 | 10/10/96 | 58 | rusty | 10 | 35.0 |
| 58 | 103 | 11/12/96 | 22 | dustin | 7 | 45.0 |
| 58 | 103 | 11/12/96 | 31 | lubber | 8 | 55.5 |
| 58 | 103 | 11/12/96 | 58 | rusty | 10 | 35.0 |



Temp1

| sid1 | bid | day | sid2 | sname | rating | age |
|------|-----|----------|------|--------|--------|------|
| 22 | 101 | 10/10/96 | 22 | dustin | 7 | 45.0 |
| 22 | 101 | 10/10/96 | 31 | lubber | 8 | 55.5 |
| 22 | 101 | 10/10/96 | 58 | rusty | 10 | 35.0 |
| 58 | 103 | 11/12/96 | 22 | dustin | 7 | 45.0 |
| 58 | 103 | 11/12/96 | 31 | lubber | 8 | 55.5 |
| 58 | 103 | 11/12/96 | 58 | rusty | 10 | 35.0 |

Relational Algebra Operators



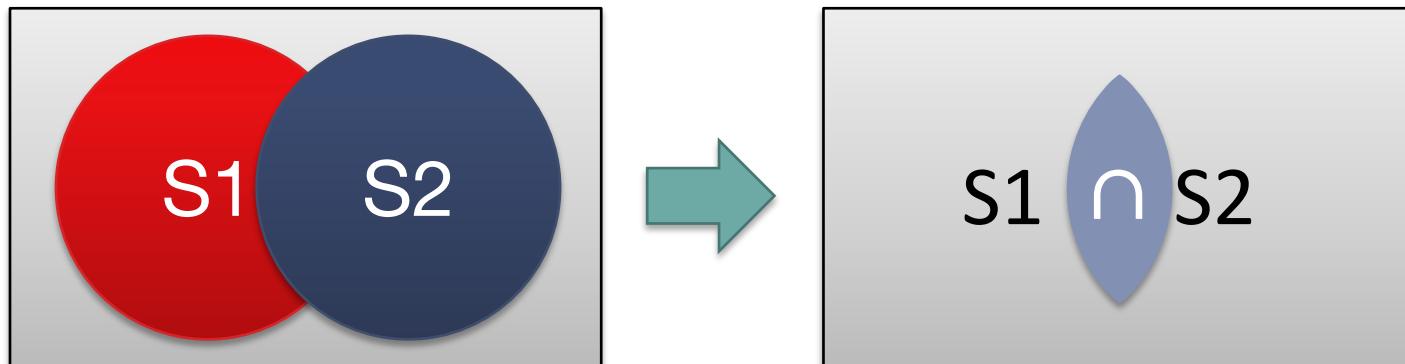
- **Projection** (π): Retains only desired columns (vertical) *select*
- **Selection** (σ): Selects a subset of rows (horizontal) *where*
- **Renaming** (ρ): Rename attributes and relations
- **Union** (\cup): Tuples in r1 or in r2.
- **Set-difference** ($-$): Tuples in r1, but not in r2.
- **Cross-product** (\times): Allows us to combine two relations.
- **Intersection** (\cap): Tuples in r1 and in r2.  *Next*
- **Joins** (\bowtie_θ , \bowtie): Combine relations that satisfy predicates

Compound Operator: Intersection

Berkeley
cs186

- Same as with union, both input relations must be *compatible*.
- SQL Expression: **INTERSECT**

$$S1 \cap S2$$



Intersection (\cap)



- Same story as $-$ with respect to duplicates
 - Duplicates don't need to be eliminated if inputs are sets

Relational Instance $S1$

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

Relational Instance $S2$

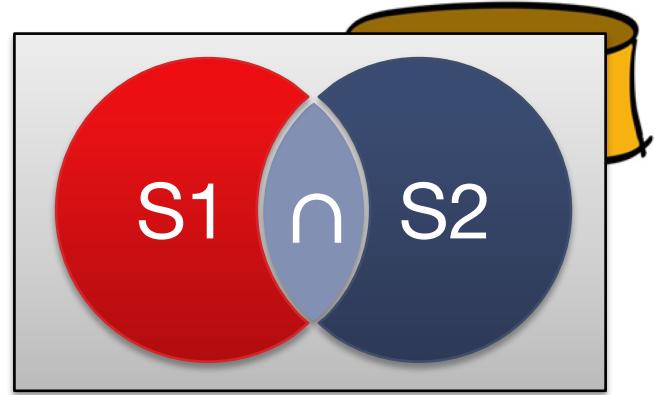
| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 28 | yuppy | 9 | 35.0 |
| 31 | lubber | 8 | 55.5 |
| 44 | guppy | 5 | 35.0 |
| 58 | rusty | 10 | 35.0 |

$S1 \cap S2$

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

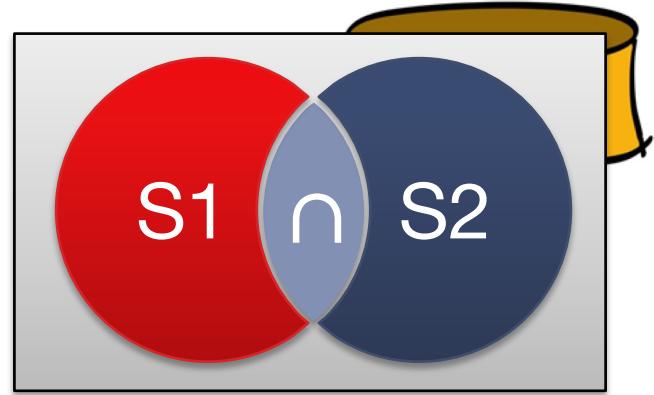
Intersection (\cap), Pt 2

- We saw that \cap is a compound operator.
- $S1 \cap S2 = ?$



Intersection (\cap), Pt 3

- $S1 \cap S2 = S1 - ?$
- Q: What is “?”

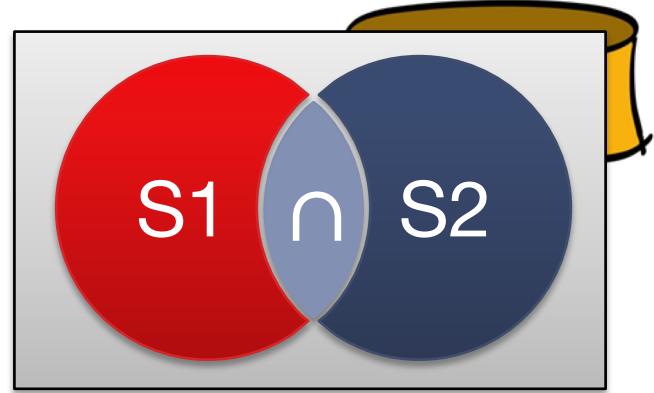


$$\text{?} = S1 - ?$$

A diagram illustrating the set difference operation. It shows a blue oval on the left followed by an equals sign. To the right of the equals sign is a red circle containing the text "S1". To the right of the circle is a minus sign. To the right of the minus sign is another red circle containing a question mark. This visualizes the concept of removing elements from S1 that are also in the set represented by the blue oval.

Intersection (\cap), Pt 4

- $S_1 \cap S_2 = S_1 - (S_1 - S_2)$



$$\text{---} = S_1 - (S_1 - S_2)$$

Diagram illustrating the set difference $S_1 - (S_1 - S_2)$. It shows three circles: a red circle labeled S_1 , a blue circle labeled S_2 , and an orange circle labeled $S_1 - S_2$. The expression $S_1 - (S_1 - S_2)$ is shown as a subtraction operation where the result is the blue circle S_2 . Below the circles, a purple hand-drawn sketch shows a irregular shape with the formula $S_1 - S_2$ written inside it, with two vertical lines above it.

Compound Operator: Join



- Joins are compound operators (like intersection):
 - Generally, $\sigma_\theta(R \times S)$
 - With possibly a rename in there (for natural join)
- Increasing degree of specialization
 - **Theta Join** (\bowtie_θ): join on logical expression θ
 - **Equi-Join**: theta join with theta being a conjunction of equalities
 - **Natural Join** (\bowtie): equi-join on all matching column names
 - Only one copy per column preserved!
- Relating information across tables using joins/cross-products is super useful and important
 - Want to avoid cross-products
 - We just learned about efficient join algorithms

Theta Join (\bowtie_θ) Semantics



- $R \bowtie_\theta S = \underline{\sigma_\theta(R \times S)}$ *R × S where σ*
- **Apply a cross-product, then filter out tuples that don't match.**
- If θ only contains equality conditions (with an AND between them), this is called an *equi-join*

Theta Join (\bowtie_{θ}) Example

- We want to find boats that people have reserved
- $R1 \bowtie_{sid=sid} S1$
 - Confusing... hard to interpret!

| R1: | | | S1: | | | | | | | | | | | |
|------------|------------|------------|------------|--------------|---------------|------------|------------|------------|------------|------------|--------------|---------------|------------|--|
| <u>sid</u> | <u>bid</u> | <u>day</u> | <u>sid</u> | <u>sname</u> | <u>rating</u> | <u>age</u> | <u>sid</u> | <u>bid</u> | <u>day</u> | <u>sid</u> | <u>sname</u> | <u>rating</u> | <u>age</u> | |
| 22 | 101 | 10/10/96 | 22 | dustin | 7 | 45.0 | | | | 22 | dustin | 7 | 45.0 | |
| 58 | 103 | 11/12/96 | 31 | lubber | 8 | 55.5 | | | | 58 | rusty | 10 | 35.0 | |
| | | | 58 | rusty | 10 | 35.0 | | | | | | | | |

$\bowtie_{sid=sid}$

=

- Q: How do we fix?

Theta Join (\bowtie_{θ}) Example



- $\rho_{(\text{sid} \rightarrow \text{sid1})}(\mathbf{R1}) \bowtie_{\text{sid1}=\text{sid}} \mathbf{S1}$

| <u>sid1</u> | <u>bid</u> | <u>day</u> |
|-------------|------------|------------|
| 22 | 101 | 10/10/96 |
| 58 | 103 | 11/12/96 |

S1:

| <u>sid</u> | <u>sname</u> | <u>rating</u> | <u>age</u> |
|------------|--------------|---------------|------------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

$\bowtie_{\text{sid1}=\text{sid}}$

=

| <u>sid1</u> | <u>bid</u> | <u>day</u> | <u>sid</u> | <u>sname</u> | <u>rating</u> | <u>age</u> |
|-------------|------------|------------|------------|--------------|---------------|------------|
| 22 | 101 | 10/10/96 | 22 | dustin | 7 | 45.0 |
| 58 | 103 | 11/12/96 | 58 | rusty | 10 | 35.0 |

Another Theta Join (\bowtie_θ) Self Join Example

- $R \bowtie_\theta S = \sigma_\theta(R \times S)$
- Example:** More senior sailors for each sailor.
- $\rho_{(sid \rightarrow sid1, sname \rightarrow sname1, rating \rightarrow rating1, age \rightarrow age1)} S1 \bowtie_{age1 < age} S1$

S1:

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

| S1 | | | | S1 | | | |
|------|--------|---------|------|-----|--------|--------|------|
| sid1 | sname1 | rating1 | age1 | sid | sname | rating | age |
| 22 | dustin | 7 | 45.0 | 22 | dustin | 7 | 45.0 |
| 22 | dustin | 7 | 45.0 | 31 | lubber | 8 | 55.5 |
| 22 | dustin | 7 | 45.0 | 58 | rusty | 10 | 35.0 |
| 31 | lubber | 8 | 55.5 | 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 | 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 | 58 | rusty | 10 | 35.0 |
| 58 | rusty | 10 | 35.0 | 58 | rusty | 10 | 35.0 |

Natural Join (\bowtie)



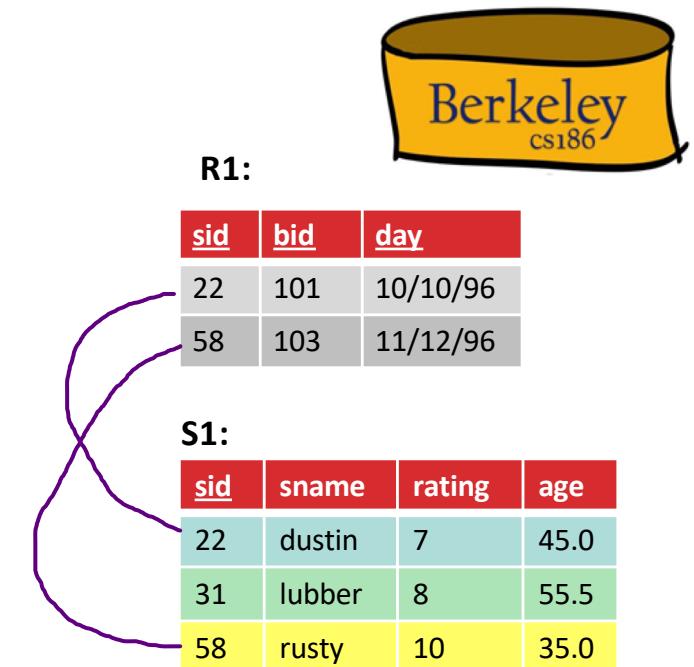
- Special case of equi-join in which equalities are specified for all matching fields and duplicate fields are projected away
- Compute $R \times S$
- Select rows where fields appearing in both relations have equal values
- Project onto the set of all unique fields.

Natural Join (\bowtie) Pt 2

- $R \bowtie S$

$R1 \bowtie S1$

| <u>sid</u> | <u>bid</u> | <u>day</u> | <u>sid</u> | <u>sname</u> | <u>rating</u> | <u>age</u> |
|------------|------------|------------|------------|--------------|---------------|------------|
| 22 | 101 | 10/10/96 | 22 | dustin | 7 | 45.0 |
| 22 | 101 | 10/10/96 | 31 | lubber | 8 | 55.5 |
| 22 | 101 | 10/10/96 | 58 | rusty | 10 | 35.0 |
| 58 | 103 | 11/12/96 | 22 | dustin | 7 | 45.0 |
| 58 | 103 | 11/12/96 | 31 | lubber | 8 | 55.5 |
| 58 | 103 | 11/12/96 | 58 | rusty | 10 | 35.0 |



Natural Join (\bowtie) Pt 3

- $R \bowtie S$

$R1 \bowtie S1$

| sid | bid | day | sname | rating | age |
|-----|-----|----------|--------|--------|------|
| 22 | 101 | 10/10/96 | dustin | 7 | 45.0 |
| 58 | 103 | 11/12/96 | rusty | 10 | 35.0 |



R1:

| sid | bid | day |
|-----|-----|----------|
| 22 | 101 | 10/10/96 |
| 58 | 103 | 11/12/96 |

S1:

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

- Commonly used for foreign key joins (as above).

Natural Join (\bowtie) Pt 3

- $R \bowtie S$

$R1 \bowtie S1$

| sid | bid | day | sname | rating | age |
|-----|-----|----------|--------|--------|------|
| 22 | 101 | 10/10/96 | dustin | 7 | 45.0 |
| 58 | 103 | 11/12/96 | rusty | 10 | 35.0 |



R1:

| sid | bid | day |
|-----|-----|----------|
| 22 | 101 | 10/10/96 |
| 58 | 103 | 11/12/96 |

S1:

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

- Q: How do we express $R1 \bowtie S1$ using the other operators?
- $R1 \bowtie S1 = \pi_{\text{sid}, \text{bid}, \text{day}, \text{sname}, \text{rating}, \text{age}} \sigma_{\text{sid} = \text{sid1}}(R1 \times \rho_{(\text{sid} \rightarrow \text{sid1})}(S1))$
- $R1 \bowtie S1 = \pi_{\text{unique fld.}} \sigma_{\text{eq. matching fld.}}(R1 \times \rho_{\text{eq. matching fld. renamed}}(S1))$



Other Join Variants

- We have convenient symbols for Outer joins:
 - **Left Outer join**
 - $R \bowtie S$
 - **Right Outer join**
 - $R \bowtie S$
 - **Full Outer join**
 - $R \bowtie S$

Complex Relational Algebra Expressions



- Algebras allow us to express sequences of operations in a natural way.
- Example
 - in arithmetic algebra: $(x + 4)^*(y - 3)$
- Relational algebra allows the same.
- Three notations:
 1. Sequences of assignment statements.
 2. Expressions with several operators.
 3. Expression trees.

Sequences of Assignments



- Create temporary relation names.
variable
- Renaming can be implied by giving relations a list of attributes.
 - $R3(\underline{X}, \underline{Y}) := R1$
take X, Y from R1
- Example: $R3 := R1 \bowtie_C R2$ can be written:
 $R4 := R1 \times R2$
 $R3 := \sigma_C(R4)$

Expressions with Several Operators



Precedence of relational operators:

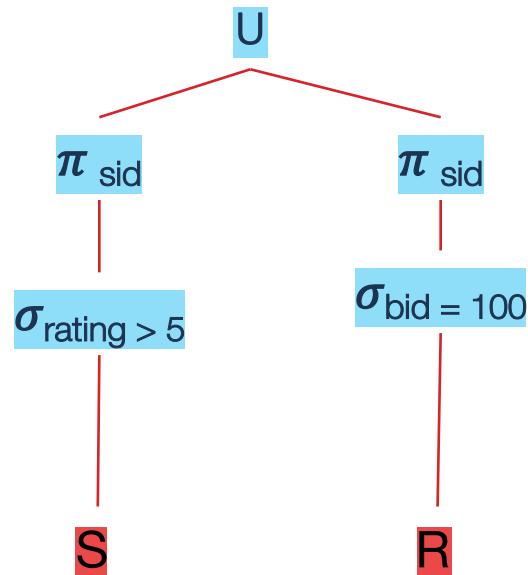
1. Unary operators --- select, project, rename --- have highest precedence, bind first.
2. Then come products and joins.
3. Then set operations bind last.

But we can always insert parentheses to force the order you desire.

Expression Trees



- Leaves are operands (relations).
- Interior nodes are operators, applied to their child or children.
- Ex: Given $R(sid, bid, day)$, $S(sid, sname, rating, age)$, find the sids of all the sailors whose rating > 5 or have reserved boat 100.



Rewriting relations in algebra is more clear than SQL

A Step Back: Why Did We Study This?



- Relational algebra expressions, just like linear algebra or elementary algebra expressions are easy to manipulate for the DBMS
- Also the number of operators is small so it's easy to work with.
- To figure out how to rewrite and simplify rel alg expressions, the DBMS uses:
 - Various heuristics
 - Various cost functions

Simple Rewritings



- Example: Changing the order of predicate evaluation
 - $\sigma_{\text{exp1} \wedge \text{exp2}} R = \sigma_{\text{exp1}} (\sigma_{\text{exp2}} R) = \sigma_{\text{exp2}} (\sigma_{\text{exp1}} R)$
- Example: Changing the order of joins
 - $(R \bowtie S) \bowtie T = R \bowtie (S \bowtie T)$

An Example of a “Rewrite”: Push-Down



- Want reservations for sailors whose age > 40

$\sigma_{\text{age} > 40} (R1 \bowtie S1)$

| sid | bid | day | sname | rating | age |
|-----|-----|----------|--------|--------|------|
| 22 | 101 | 10/10/96 | dustin | 7 | 45.0 |
| 58 | 103 | 11/12/96 | rusty | 10 | 35.0 |

| sid | bid | day |
|-----|-----|----------|
| 22 | 101 | 10/10/96 |
| 58 | 103 | 11/12/96 |

S1:

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

Q: Any other expressions?

cheaper

- Another equivalent expr: $R1 \bowtie \sigma_{\text{age} > 40} S1$
- This may be cheaper to compute!

An Example of a “Rewrite”: Eliminating Nesting



- Names of sailors who've not reserved boat #103:

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sid NOT IN
       (SELECT R.sid
        FROM   Reserves R
        WHERE  R.bid=103)
```

R:

| sid | bid | day |
|-----|-----|----------|
| 22 | 101 | 10/10/96 |
| 58 | 103 | 11/12/96 |

S:

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

An Example of a “Rewrite”: Eliminating Nesting



- Names of sailors who've not reserved boat #103:

One approach to remove inner query:

$$\pi_{\text{sname}}(R) - \pi_{\text{sname}}((\sigma_{\text{bid}=103} R) \bowtie S))$$

R:

| sid | bid | day |
|-----|-----|----------|
| 22 | 101 | 10/10/96 |
| 58 | 103 | 11/12/96 |

S:

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 31 | lubber | 8 | 55.5 |
| 58 | rusty | 10 | 35.0 |

Extended Relational Algebra



- Group By / Aggregation Operator (γ):
 - $\gamma_{age, AVG(rating)}(Sailors)$
 - With selection (HAVING clause):
 - $\gamma_{age, AVG(rating), COUNT(*)>2}(Sailors)$
- Implicitly combines GROUP BY, HAVING and SELECT

Summary



- Relational Algebra: a small set of operators mapping relations to relations
 - Operational, in the sense that you specify the explicit order of operations
 - A closed set of operators! Mix and match.
 - Easy to manipulate/rewrite/simplify
 - Super powerful! Can encapsulate a lot of SQL functionality
- Basic ops include: σ , π , \times , U , $-$
- Important compound ops: \cap , \bowtie