

Transactions & Concurrency Control I

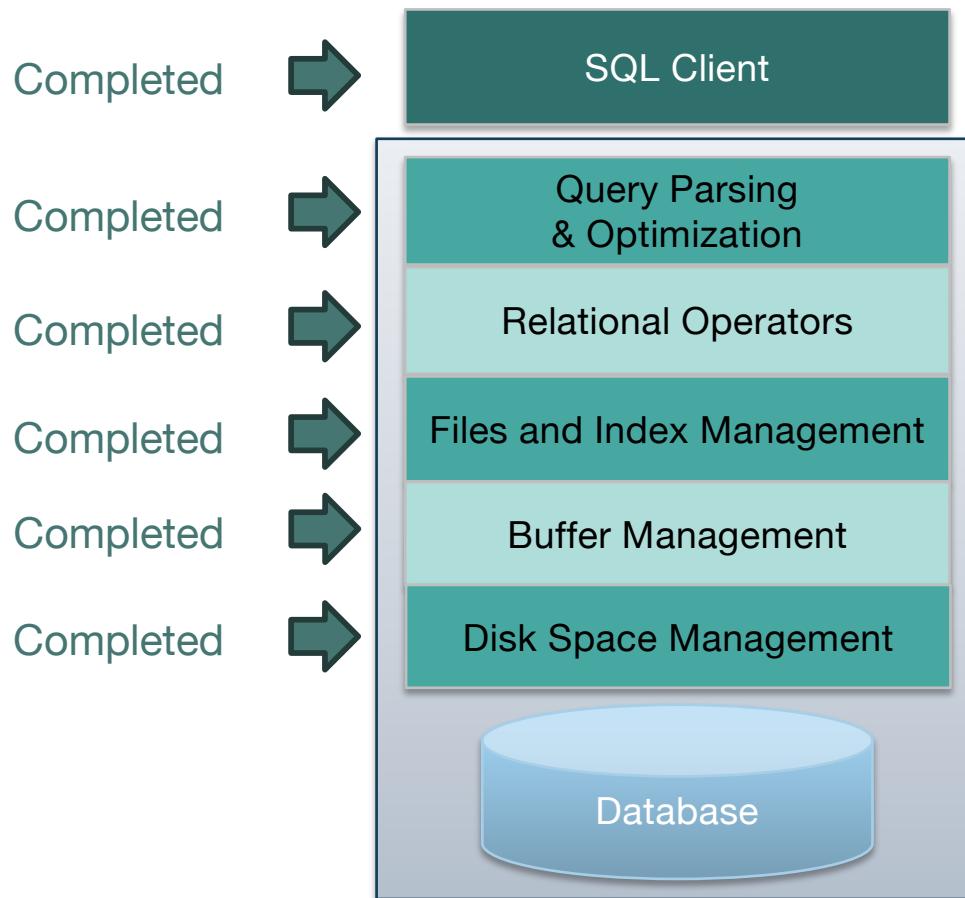
Alvin Cheung

Fall 2022

Reading: R&G Ch 16,17



Architecture of a DBMS

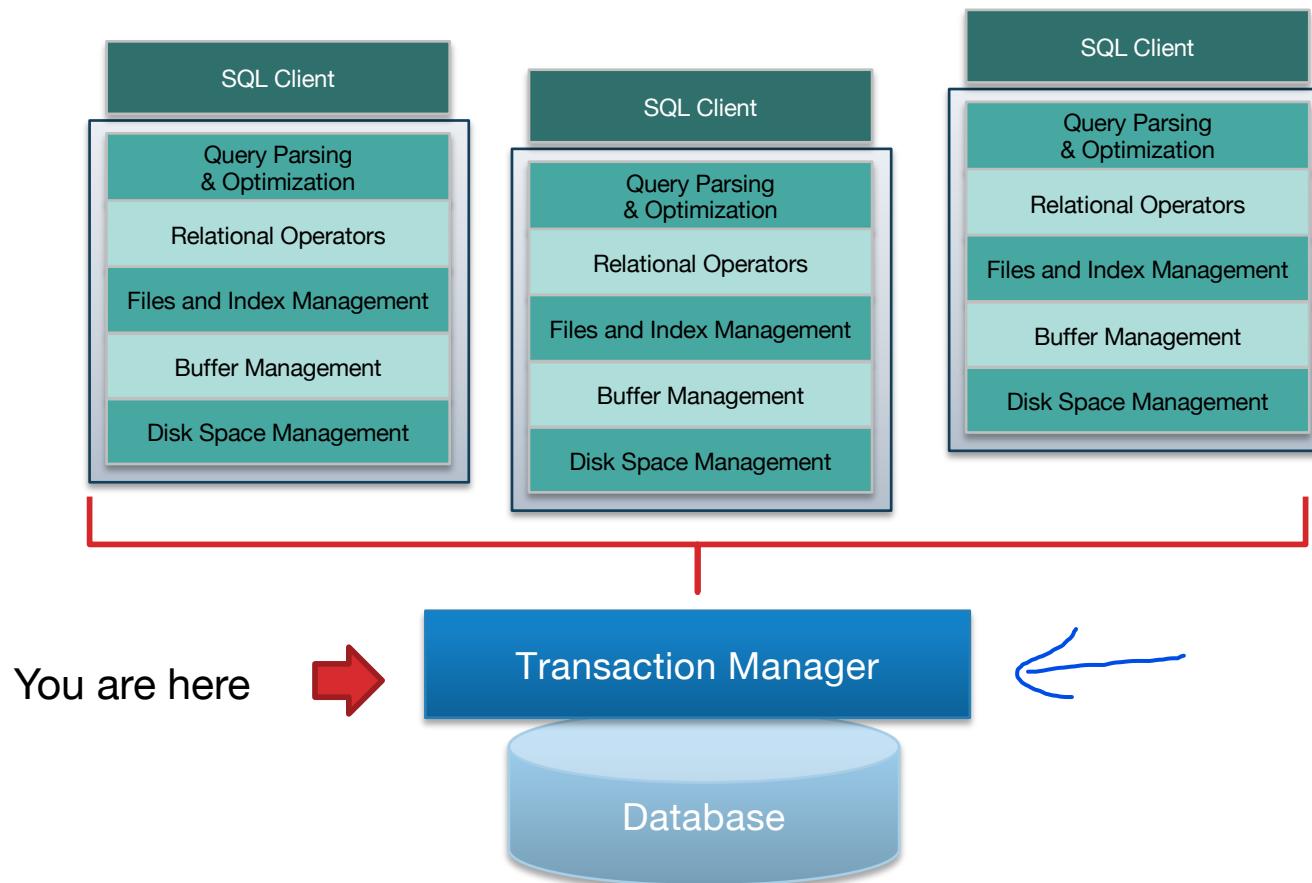


We've finished our DBMS end-to-end stack!

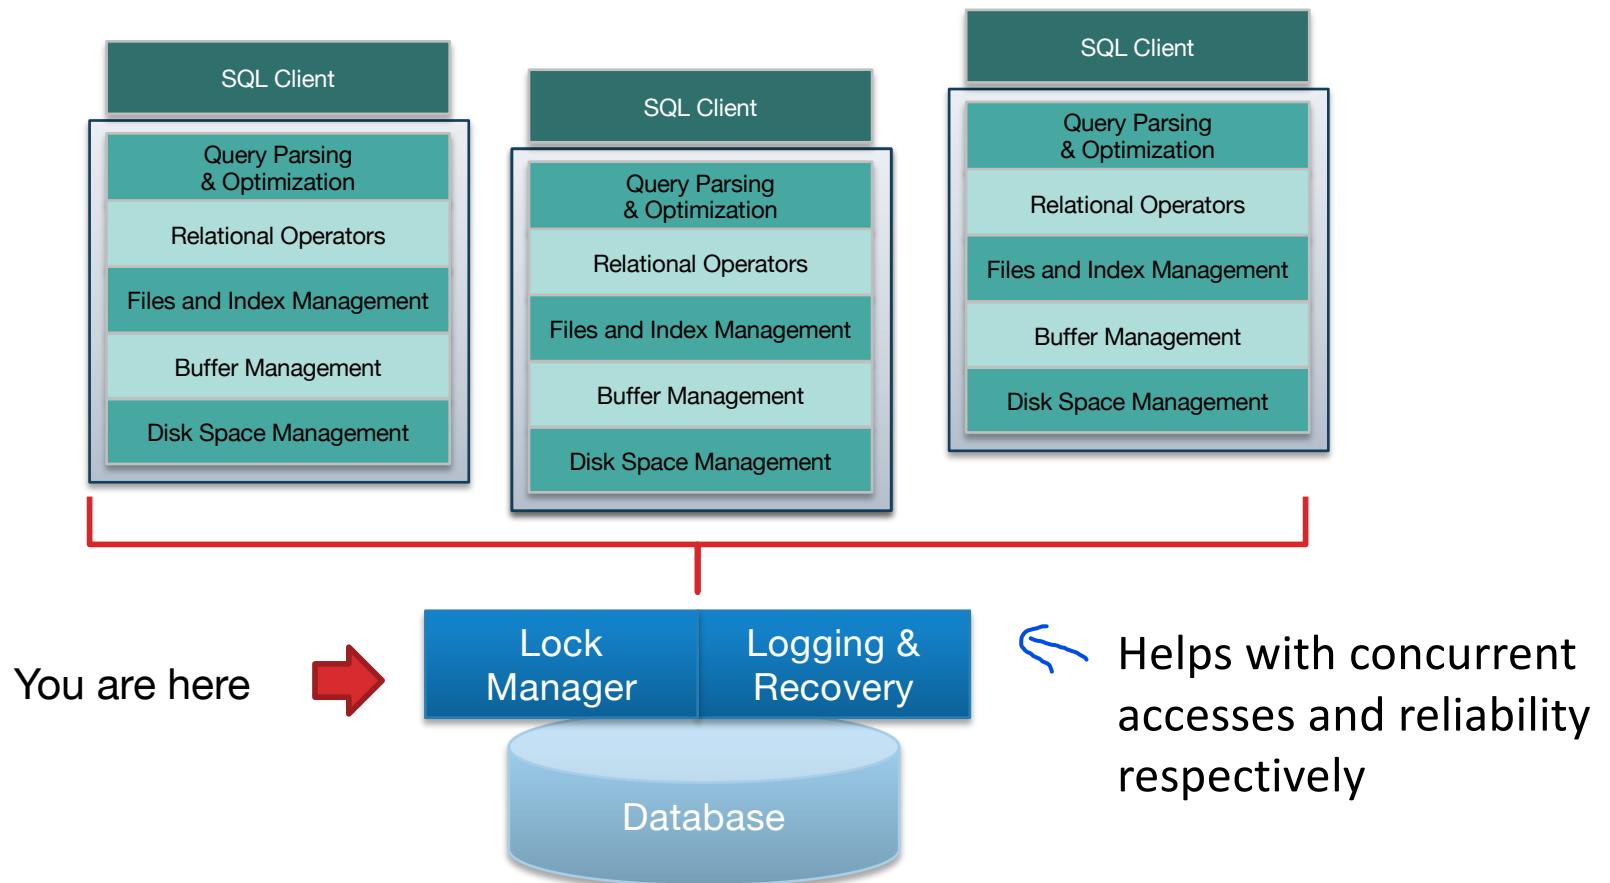
But, can't support multiple user accesses or be reliable to failures.

That is the focus of today's lecture.

Architecture of a DBMS, Part 2



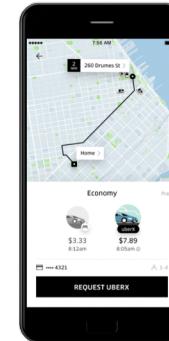
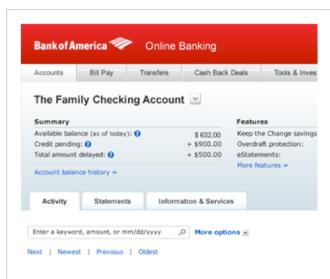
Architecture of a DBMS, Part 3



Applications on DBMS



- Virtually any service that maintains state today is an application on top of some kind of DBMS
 - Uber/Lyft
 - Kayak
 - Amazon
 - Bank of America
 - Tiktok

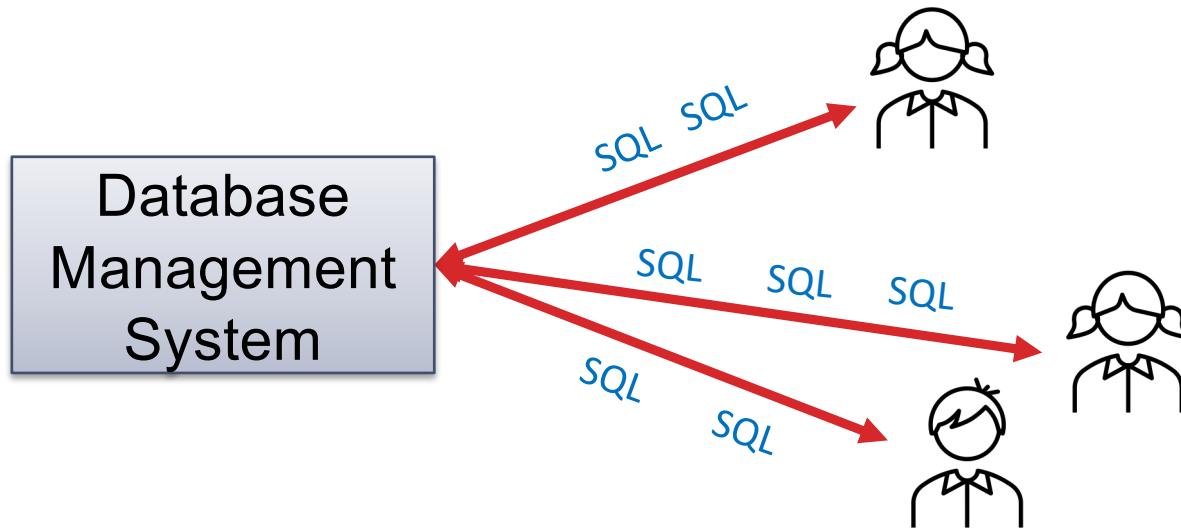


KAYAK

Applications Want Something from the DBMS



- Single queries and updates of course: what you learned so far!
- Real applications are composed of many statements being generated by user behaviors
 - E.g., transfer money from Checking to Savings as two separate SQL statements
- Many users work with the application at the same time



Concurrency Control & Recovery



- **Part 1: Concurrency Control**
 - Correct/fast data access in the presence of concurrent work by many users
 - Interleaved processing that to a user, provides the illusion of no interference
- **Part 2: Recovery**
 - Ensure database is fault tolerant
 - Not corrupted by application, DBMS, power, or media failure
 - Storage guarantees for mission-critical data
- **It's all about the programmer!**
 - Systems provide guarantees
 - These guarantees lighten the load of app writers

Concurrent Execution: Why bother?



- Multiple queries are allowed to run concurrently in the system.
- Advantages are twofold:
 - Throughput (queries per second):
 - Increase processor/disk utilization
 - Single core: one query can use the CPU while another is reading to/writing from the disk
 - Multicore: ideally, scale throughput in the number of processors
 - Latency (response time per query):
 - Multiple queries can run at the same time rather than waiting for earlier ones to finish
 - So one query's latency need not be dependent on another unrelated one's
 - Lightweight queries are not bottlenecked on more time-consuming ones to finish
 - Or that's the hope
 - Both are important!

What can go wrong?

Ex |



Moving budget from advertising to inventory and sales

```
UPDATE Budget
```

```
SET money = money - 500
```

```
WHERE category = "advertising"
```

```
SELECT SUM(money)  
FROM Budget
```

```
UPDATE Budget
```

```
SET money = money + 200
```

```
WHERE category = "inventory"
```

```
UPDATE Budget
```

```
SET money = money + 300
```

```
WHERE category = "salaries"
```

Money not conserved!

Two Issues:

1. Order of operations matters!
2. Users need a way to say what's acceptable
e.g., can RHS query view the budget in between the update statements?

This is called an inconsistent read
aka a **WRITE-READ** conflict

What can go wrong? Ex 2



- App 1:

```
SELECT inventory FROM products  
WHERE pid = 1
```
- App 2:

```
UPDATE products SET inventory = 0  
WHERE pid = 1
```
- App 1:

```
SELECT inventory * price FROM products  
WHERE pid = 1
```
- This is known as an unrepeatable read
aka a **READ-WRITE** conflict

What can go wrong?

Ex 3



User 1

```
UPDATE Account  
SET amount = 10000000  
WHERE number = "my-account"
```

User 2

```
SELECT amount  
FROM Account  
WHERE number = "my-account"
```

Aborted by
the system

What if User 2 read the amount of my-account before User 1 aborted?

And then proceeded to buy a house, car, and yacht!

What could go wrong? This is called the dirty read aka **WRITE-READ** conflict

What can go wrong?

Ex 4

Account 1 = \$100
Account 2 = \$100
Total = \$200



- App 1:
 - Set Account 1 = \$200
 - Set Account 2 = \$0
- App 2:
 - Set Account 2 = \$200
 - Set Account 1 = \$0
- At the end:
 - Total = \$200
- App 1: Set Account 1 = \$200
- App 2: Set Account 2 = \$200
- App 1: Set Account 2 = \$0
- App 2: Set Account 1 = \$0
- At the end:
 - Total = \$0

timeline

This is called the lost update aka a WRITE-WRITE conflict

What can go wrong?

- Buying tickets to the next Bieber concert:
 - Fill up form with your mailing address
 - Put in debit card number
 - Click submit
 - Screen shows money deducted from your account
 - [Your browser crashes]

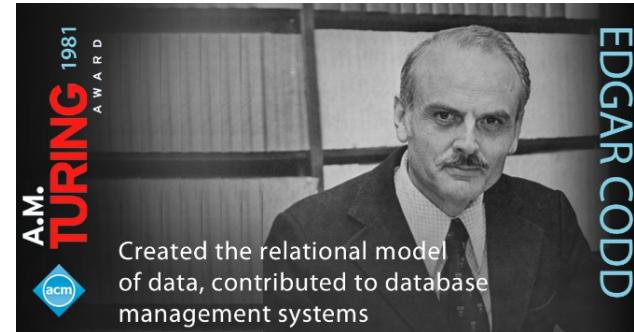
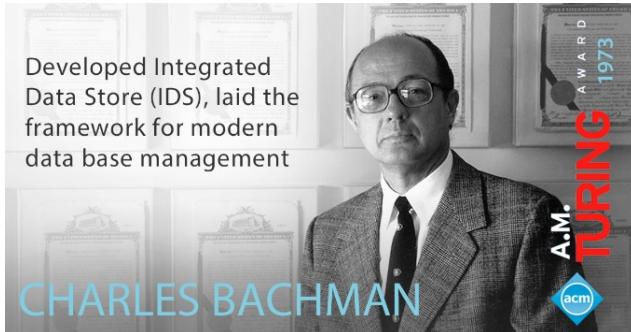


Lesson:
Changes to the database
should be **ALL or NOTHING**

Avoid problems mentioned above

TRANSACTIONS

An Aside: Database Turing Awards



What is a Transaction?



- A sequence of *multiple actions* to be executed as an atomic unit
 - All of it, or none of it are reflected on the database
- Application View (SQL View):
 - Begin transaction
 - Sequence of SQL statements
 - End transaction
- Examples
 - Transfer money between accounts
 - Book a flight, a hotel and a car together on Expedia

Our Transaction Model



- **Transaction** (“Txn/Xact”):
 - DBMS’s abstract view of an application program (or activity)
 - A sequence of *reads* and *writes* of database objects
 - Batch of work that must *commit* or *abort* as an atomic unit
- **Txn Manager controls execution of txns**
- **Program logic is invisible to DBMS!**
 - Arbitrary computation possible on data fetched from the DB
 - The DBMS only sees data read/written from/to the DB

ACID: High-Level Properties of Transactions



- **A** tomicity: All actions in the txn happen, or none happen.
- **C** onsistency: If the DB starts out consistent, it ends up consistent at the end of the txn
- **I** solation: Execution of each txn is *isolated* from that of others
- **D** urability: If a txn commits, its effects persist.
 ↑
 Looks like running one transaction at a time

Note: This is a mnemonic, not a formalism. We'll do some formalisms shortly.

Isolation (Concurrency)

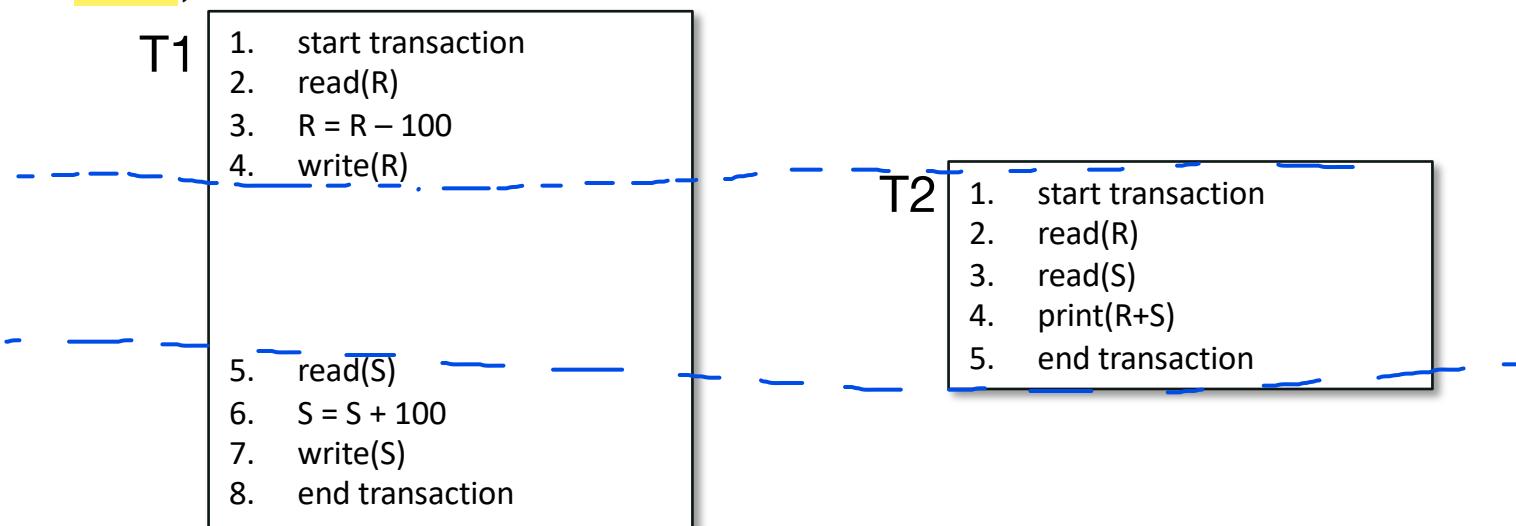


- DBMS interleaves actions of many txns
 - Actions = reads/writes of DB objects
- DBMS ensures 2 txns do not “interfere”
- Each txn executes as if it ran by itself.
 - Concurrent accesses have no effect on txn’s behavior
 - Net effect must be identical to executing all transactions in some serial order
 - Users & programmers think about transactions in isolation
 - Without considering effects of other concurrent txn!

Isolation: An Example



- Concurrency introduces problems
 - If another transaction T2 accesses R and S between steps 4 and 5 of T1, it will see a lower value for R+S.



- T2 sees state “in-between” T1’s changes. Instead want it to run before T1 entirely or after T1 entirely
- Isolation easy to achieve by running one txn at a time
 - However, recall that serial execution is not desirable

Atomicity and Durability



- A transaction ends in one of two ways:
 - Commit after completing all its actions *execute all*
 - “commit” is a contract between the app and the DBMS
 - The changes of the transactions need to be reflected in the database
 - Abort (requested by app or be aborted by the DBMS) after executing some actions *execute none*
 - Or system crash while the txn is in progress; treat as abort.
- Two key properties for a transaction
 - Atomicity: Either execute all its actions (committed), or none of them (aborted)
 - Durability: The effects of a committed txn must survive failures.
- DBMS typically ensures the above by logging all actions:
 - Undo the actions of aborted/failed transactions.
 - Redo actions of committed transactions not yet propagated to disk when system crashes

Atomicity and Durability, Example cont.



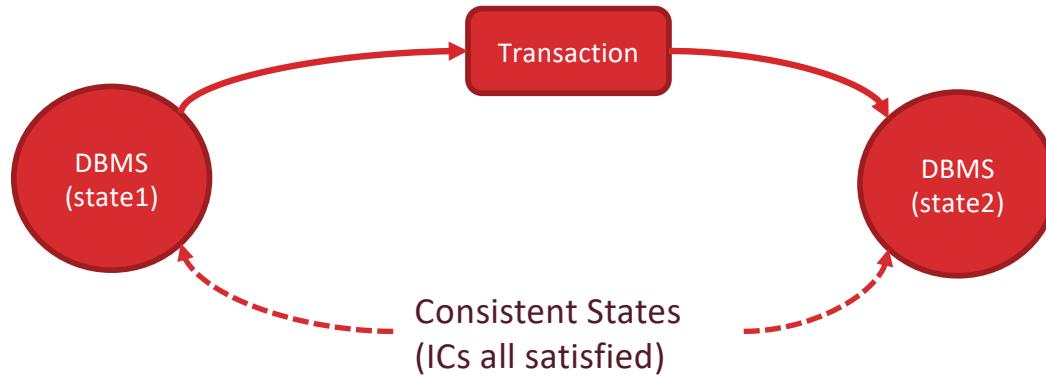
- **Atomicity**
 - If the transaction fails after step 4 and before step 7
 - Money will be “lost” → inconsistent database
 - DBMS should ensure that updates of a partially executed transaction are not reflected
- **Durability** “Effects persist”
 - Once the user hears that the transaction is complete, can rest easy that \$100M was indeed transferred from R to S.

1. start transaction
2. read(R)
3. $R = R - 100$
4. write(R)
5. read(S)
6. $S = S + 100$
7. write(S)
8. end transaction

Transaction Consistency



- **Transactions preserve DB consistency**
 - Given a consistent DB state, produce another consistent DB state
- DB consistency expressed as a set of **declarative integrity constraints**
 - Primary / Foreign key constraints
- **Transactions that violate integrity are aborted**



Summary



- We have seen an overview
- ACID Transactions make guarantees that
 - Improve performance (via concurrency)
 - Relieve programmers of correctness concerns
 - Hide concurrency and failure handling!
- Two key issues to consider, and mechanisms
 - Concurrency control (via two-phase locking)
 - Recovery (via write-ahead logging WAL)
- We'll do concurrency control first

CONCURRENCY CONTROL

Concurrency Control: Providing Isolation



- **Naïve approach - serial execution, no concurrency**
 - One transaction runs at a time
 - Safe but slow
- **Execution must be interleaved for better performance**
- With concurrent executions, **how does one define and ensure correctness?** What sequencing are permitted?

Transaction Schedules



Tabular representation

T1	T2
begin	
read(A)	
write(A)	
read(B)	
write(B)	
commit	
	begin
	read(A)
	write(A)
	read(B)
	write(B)
	commit

A **schedule** is a sequence of actions on data from one or more transactions.

Actions: Begin, Read, Write, Commit and Abort.

String representation:

R₁(A) W₁(A) R₁(B) W₁(B) R₂(A) W₂(A) R₂(B) W₂(B)

By convention we only include committed transactions, and **omit Begin and Commit in string representation**

Serial Equivalence

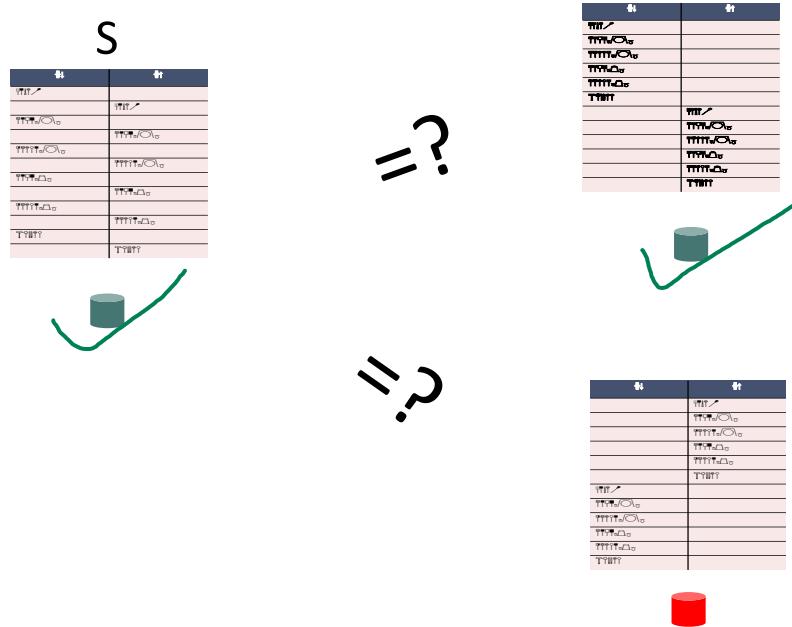


- We need a starting point for correct behavior
 - **Definition:** **Serial schedule** *Serial*
 - Each transaction runs from start to finish without any intervening actions from other transactions
 - Complete isolation
 - **Definition:** 2 schedules are **equivalent** if they:
 - involve the same transactions
 - each individual transaction's actions are ordered the same
 - both schedules leave the DB in the same final state

Serializability



- **Definition:** Schedule S is **serializable** if:
 - S is equivalent to **some** serial schedule



Schedule 1



T1: Transfer \$100 from A to B	T2: Add 10% interest to A & B
begin	
read(A)	
A = A - 100	
write(A)	
read(B)	
B = B + 100	
write(B)	
commit	
	begin
	read(A)
	A = A * 1.1
	write(A)
	read(B)
	B = B * 1.1
	write(B)
	commit

- Let T1 transfer \$100 from A to B
- Let T2 add 10% interest to A & B
- Serial schedule in which T1 is followed by T2
 - Final outcome:
 - $A := 1.1 * (A - 100)$
 - $B := 1.1 * (B + 100)$

Schedule 2



T1: Transfer \$100 from A to B	T2: Add 10% interest to A & B
	begin
	read(A)
	$A = A * 1.1$
	write(A)
	read(B)
	$B = B * 1.1$
	write(B)
	commit
begin	
read(A)	
$A = A - 100$	
write(A)	
read(B)	
$B = B + 100$	
write(B)	
commit	

- Serial schedule in which T2 is followed by T1
 - Final outcome:
 - $A := (1.1*A)-100$
 - $B := (1.1*B)+100$
 - Different!
 - But still understandable

Schedule 3



T1: Transfer \$100 from A to B	T2: Add 10% interest to A & B
begin	
read(A)	
A = A - 100	
write(A)	
	begin
	read(A)
	A = A * 1.1
	write(A)
read(B)	
B = B + 100	
write(B)	
commit	
	read(B)
	B = B * 1.1
	write(B)
	commit

- Schedule in which actions of T1 and T2 are interleaved.
- **This is not a serial schedule**
- **But it is equivalent to schedule 1**
 - $A := 1.1*(A-100)$
 - $B := 1.1*(B+100)$
- **Hence serializable!**

Conflicting Operations



- Tricky to check property “leaves the DB in the same final state”
 - Can’t really try all operation permutations and see if equivalent to a serial schedule
- Use notion of “conflicting” operations (read/write)
- Definition: Two operations conflict if they:
 - Are by different transactions,
 - Are on the same object, - At least one of them is a write.

And And
- Convince yourself: the order of non-conflicting operations has no effect on the final state of the database!
 - Focus our attention on the order of conflicting operations

$r(A) \dots w(A)$
 $w(A) \dots w(A)$

Conflict Serializable Schedules



- **Definition:** Two schedules are **conflict equivalent** if:
 - They involve the same actions of the same transactions in same order, and
 - Every pair of conflicting actions is ordered the same way
- **Definition:** Schedule **S** is **conflict serializable** if:
 - S is conflict equivalent to some serial schedule
 - Which implies that S is also Serializable [i.e., S is equivalent to a serial schedule]

$S \rightarrow S$, $S \not\rightarrow CS$

Note: some serializable schedules are NOT conflict serializable

- Conflict serializability gives false negatives as a test for serializability!
- This is the cost of a conservative test
- A price we pay to achieve efficient enforcement

Conflict Serializability - Intuition



- [Equivalent definition] A schedule S is conflict serializable if
 - We are able to transform S into a serial schedule by swapping consecutive non-conflicting operations of different transactions
- *Example*

R(A) W(A)

R(B) W(B)

R(A) W(A)

R(B) W(B)

Txn 1

Txn 2

Conflict Serializability – Intuition, Part 2



- A schedule S is conflict serializable if
 - We are able to transform S into a serial schedule by swapping consecutive non-conflicting operations of different transactions
- *Example*

R(A) W(A)	R(B) W(B)
R(A) W(A)	R(B) W(B)

R(A) W(A)	R(B) W(B)
R(A) W(A)	R(B) W(B)



Conflict Serializability – Intuition, Part 3



- A schedule S is conflict serializable if
 - We are able to transform S into a serial schedule by swapping consecutive non-conflicting operations of different transactions
- *Example*

R(A) W(A)	R(B) W(B)
R(A) W(A)	R(B) W(B)

R(A) W(A) R(B) W(B)
R(A) W(A) R(B) W(B)



Conflict Serializability – Intuition, Part 4



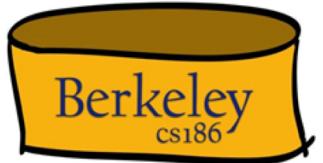
- A schedule **S** is conflict serializable if
 - We are able to transform S into a serial schedule by swapping **consecutive non-conflicting** operations of different transactions
- *Example*

R(A) W(A)	R(B) W(B)
R(A) W(A)	R(B) W(B)

R(A) W(A)	R(B) W(B)
R(A)	W(A) R(B) W(B)

3

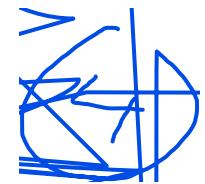
Conflict Serializability – Intuition, Part 5



- A schedule S is conflict serializable if
 - We are able to transform S into a serial schedule by swapping consecutive non-conflicting operations of different transactions
- *Example*

R(A) W(A)	R(B) W(B)
R(A) W(A)	R(B) W(B)

R(A) W(A) R(B)	W(B)
R(A)	W(A) R(B) W(B)



Conflict Serializability – Intuition, cont

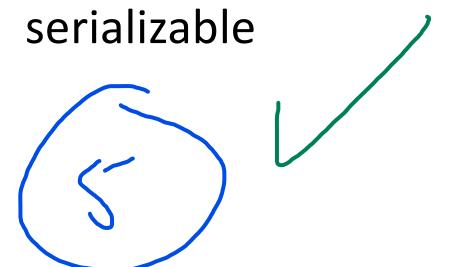


- A schedule **S** is conflict serializable if
 - We are able to transform S into a serial schedule by swapping **consecutive non-conflicting** operations of different transactions
- *Example*

R(A) W(A)	R(B) W(B)
R(A) W(A)	R(B) W(B)

R(A) W(A) R(B) W(B)
R(A)W(A) R(B) W(B)

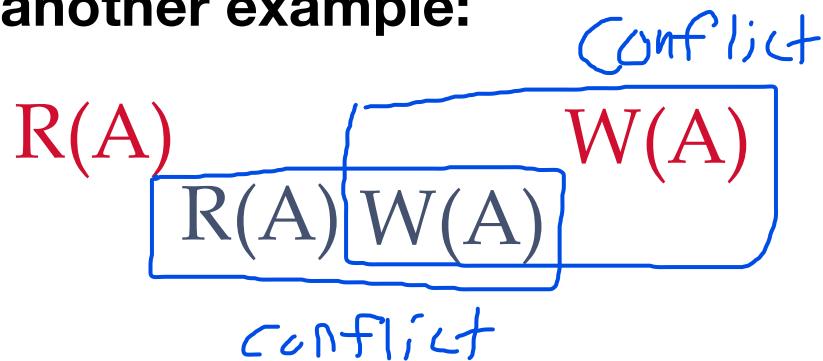
This schedule is
conflict equiv. to
a serial schedule
==> it is conflict
serializable



Conflict Serializability (Continued)



- Here's another example:



Can only swap
order of $R(A)$ s, but
nothing more.

- Q: Conflict Serializable or not?

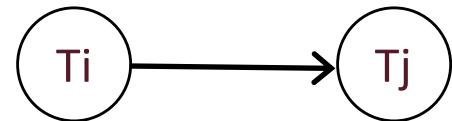
NOT!

Checking: Conflict Dependency Graph



- **Dependency Graph:**

- One node per txn
- Edge from T_i to T_j if:
 - An operation O_i of T_i conflicts with an operation O_j of T_j and
 - O_i appears earlier in the schedule than O_j



- **Theorem: Schedule is conflict serializable if and only if its dependency graph is acyclic.**

Proof Sketch: Conflicting operations prevent us from “swapping” operations into a serial schedule

Conflict Dependency Graph

Example



- A schedule that is not conflict serializable

T1: R(A), W(A)

T1

T2

Dependency graph

Example, pt 2



- A schedule that is ~~not~~ conflict serializable
(it is already serial)

T1: R(A), W(A),

T2: R(A)



Example, pt 3



- A schedule that is ~~not~~ conflict serializable

T1: R(A), W(A),

T2: R(A), W(A), R(B), W(B)



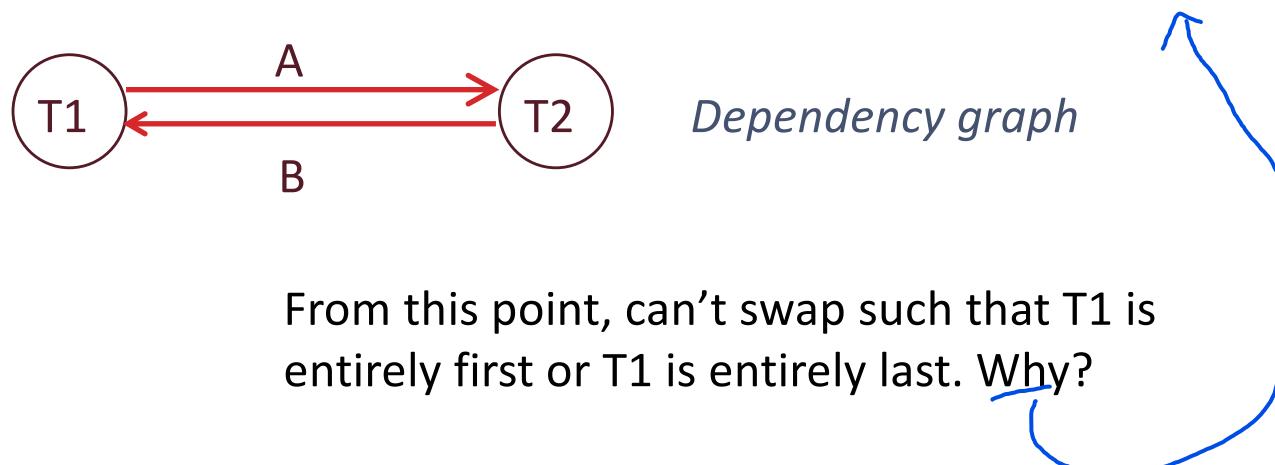
Dependency graph

Example, pt 4



- A schedule that is not conflict serializable

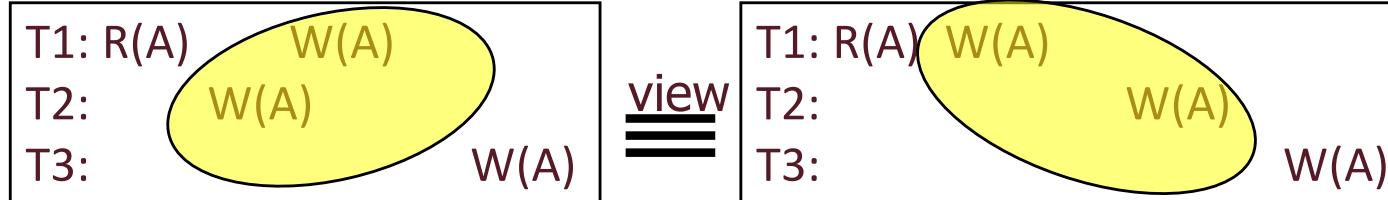
T1:	R(A), W(A),
T2:	R(A), W(A), R(B), W(B)



View Serializability



- Alternative notion of serializability: fewer false negatives
- Schedules S_1 and S_2 are view equivalent if:
 - Same initial reads:
 - If T_i reads initial value of A in S_1 , then T_i also reads initial value of A in S_2
 - Same dependent reads:
 - If T_i reads value of A written by T_j in S_1 , then T_i also reads value of A written by T_j in S_2
 - Same winning writes:
 - If T_i writes final value of A in S_1 , then T_i also writes final value of A in S_2
- Basically, allows all conflict serializable schedules + “blind writes”



Non-conflicting = everything
except R(A)W(A)

don't matter b/c W(A) at end overwrite

Notes on Serializability Definitions



- **View Serializability allows (a few) more schedules than conflict serializability**
 - But view serializability is difficult to enforce efficiently.
- **Neither definition allows all schedules that are actually serializable.**
 - Because they don't understand the meanings of the operations or the data
- **Conflict Serializability is what gets used, because it can be enforced efficiently**

↑
View Serializability is NP-complete