

# **Transactions & Concurrency Control II**

Alvin Cheung  
Fall 2022

Reading: R & G Chapter 16-17



# Transaction Implementations



- Many available!
  - Targets different workloads
- We will focus on **lock-based** implementations
  - Others: use multiple versions of data and “optimistically” let transactions move forward
    - Abort when conflicts are detected
  - Some names to know/look up:
    - Optimistic Concurrency Control
    - Timestamp-Ordered Multiversion Concurrency Control
  - We will not study these schemes in this lecture

# “Lock” data??

- Not by any crypto or hardware enforcement
  - There are no adversaries here ... this is all within the DBMS
- Recall locks / semaphores from 61c
  - These are synchronization primitives
  - Locking / unlocking has costs
- We lock by simple convention within the DBMS:
  - Each *data element* has a unique *lock*  $L_1, L_2, \dots$
  - Each transaction must first acquire the lock before reading/writing that element
  - If the lock is taken by another transaction, then wait
  - The transaction must release the lock(s) at some point
- Different *lock protocols / schemes* differ by:
  - When to lock / unlock each data element
  - What data element to lock
  - What happens when a txn waits for a lock



# What are “data elements”?



Major differences between vendors:

- Lock on the entire database
  - SQLite
- Lock on individual records
  - SQL Server, DB2, etc
- Will see tradeoffs later on

# Actions on Locks



$\text{Lock}_i(A) / L_i(A)$  = transaction  $T_i$  acquires lock for element A

$\text{Unlock}_i(A) / U_i(A)$  = transaction  $T_i$  releases lock for element A

Let's see this in action...

# A Non-S Serializable Schedule



T1

READ(A)  
A := A+100  
WRITE(A)

READ(B)  
B := B+100  
WRITE(B)

T2

READ(A)  
A := A\*2  
WRITE(A)  
READ(B)  
B := B\*2  
WRITE(B)



# Example

T1

$L_1(A)$ ; READ(A)  
A := A+100  
WRITE(A);  $U_1(A)$ ;  $L_1(B)$

READ(B)  
B := B+100  
WRITE(B);  $U_1(B)$ ;

T2

$L_2(A)$ ; READ(A)  
A := A\*2  
WRITE(A);  $U_2(A)$ ;  
 $L_2(B)$ ; BLOCKED...

...GRANTED; READ(B)  
B := B\*2  
WRITE(B);  $U_2(B)$ ;

Using locks has ensured a conflict-serializable schedule



# But...



T1

$L_1(A)$ ; READ(A)  
A := A+100  
WRITE(A);  $U_1(A)$ ;

T2

$L_2(A)$ ; READ(A)  
A := A\*2  
WRITE(A);  $U_2(A)$ ;  
 $L_2(B)$ ; READ(B)  
B := B\*2  
WRITE(B);  $U_2(B)$ ;



$L_1(B)$ ; READ(B)  
B := B+100  
WRITE(B);  $U_1(B)$ ;

Locks did not enforce conflict-serializability!!! What's wrong ?

# Two Phase Locking (2PL)



The 2PL rule:

In every transaction, all lock requests must precede all unlock requests

*before*

# Example: 2PL transactions



T1

L<sub>1</sub>(A); L<sub>1</sub>(B); READ(A)  
A := A+100  
WRITE(A); U<sub>1</sub>(A)

READ(B)  
B := B+100  
WRITE(B); U<sub>1</sub>(B)

Now it is conflict-serializable

T2

L<sub>2</sub>(A); READ(A)  
A := A\*2  
WRITE(A);  
L<sub>2</sub>(B); BLOCKED...

...GRANTED; READ(B)

B := B\*2  
WRITE(B); U<sub>2</sub>(A); U<sub>2</sub>(B)

all locks precede unlocks

# Two Phase Locking (2PL)



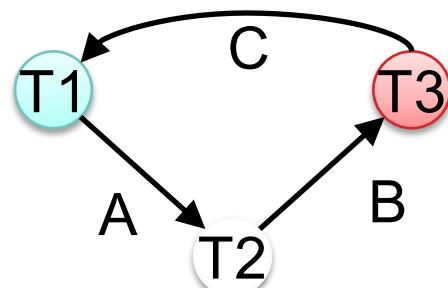
**Theorem:** 2PL ensures conflict serializability

# Two Phase Locking (2PL)



**Theorem:** 2PL ensures conflict serializability

**Proof.** Suppose not: then  
there exists a cycle  
in the dependence graph.



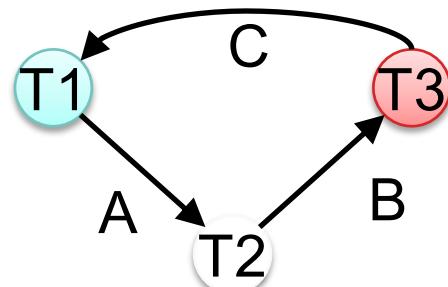
# Two Phase Locking (2PL)



**Theorem:** 2PL ensures conflict serializability

**Proof.** Suppose not: then there exists a cycle in the dependence graph.

Then there is the following temporal cycle in the schedule:

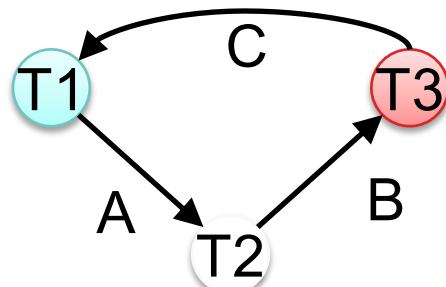


# Two Phase Locking (2PL)



**Theorem:** 2PL ensures conflict serializability

**Proof.** Suppose not: then there exists a cycle in the dependence graph.



Then there is the following temporal cycle in the schedule:

$U_1(A) \rightarrow L_2(A)$  why?

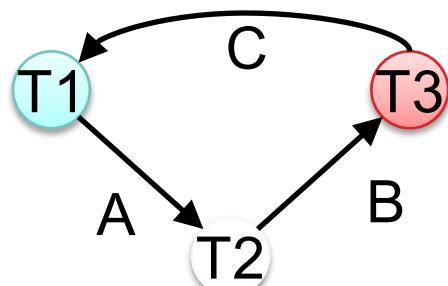
$U_1(A)$  happened strictly before  $L_2(A)$

# Two Phase Locking (2PL)



**Theorem:** 2PL ensures conflict serializability

**Proof.** Suppose not: then there exists a cycle in the dependence graph.



Then there is the following temporal cycle in the schedule:

$U_1(A) \rightarrow L_2(A)$

$L_2(A) \rightarrow U_2(B)$

why?

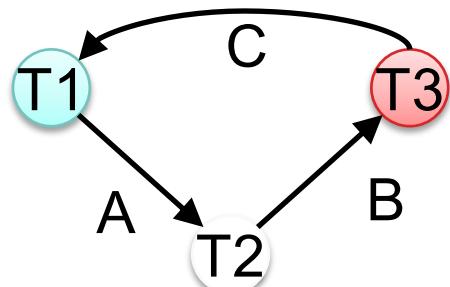
$L_2(A)$  happened strictly before  $U_1(A)$

# Two Phase Locking (2PL)



**Theorem:** 2PL ensures conflict serializability

**Proof.** Suppose not: then there exists a cycle in the dependence graph.



Then there is the following temporal cycle in the schedule:

$$U_1(A) \rightarrow L_2(A)$$

$$L_2(A) \rightarrow U_2(B)$$

$$U_2(B) \rightarrow L_3(B)$$

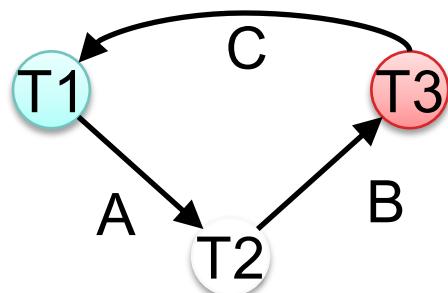
why?

# Two Phase Locking (2PL)



**Theorem:** 2PL ensures conflict serializability

**Proof.** Suppose not: then there exists a cycle in the dependence graph.



Then there is the following temporal cycle in the schedule:

$U_1(A) \rightarrow L_2(A)$

$L_2(A) \rightarrow U_2(B)$

$U_2(B) \rightarrow L_3(B)$

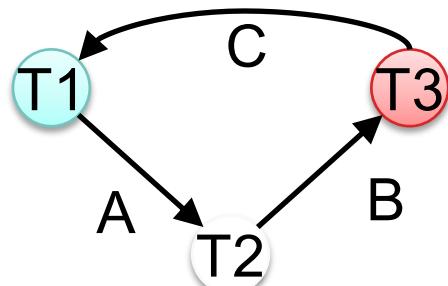
.....etc.....

# Two Phase Locking (2PL)



**Theorem:** 2PL ensures conflict serializability

**Proof.** Suppose not: then there exists a cycle in the dependence graph.



Then there is the following temporal cycle in the schedule:

- $U_1(A) \rightarrow L_2(A)$
- $L_2(A) \rightarrow U_2(B)$
- $U_2(B) \rightarrow L_3(B)$
- $L_3(B) \rightarrow U_3(C)$
- $U_3(C) \rightarrow L_1(C)$
- $L_1(C) \rightarrow U_1(A)$

Cycle in time:  
Contradiction

# A New Problem: Non-recoverable Schedule



T1

$L_1(A); L_1(B);$  READ(A)

$A := A + 100$

WRITE(A);  $U_1(A)$

READ(B)

$B := B + 100$

WRITE(B);  $U_1(B)$ ;

Rollback *(abort)*

T2

$L_2(A);$  READ(A)

$A := A * 2$

WRITE(A);

$L_2(B);$  BLOCKED...

...GRANTED; READ(B)

$B := B * 2$

WRITE(B);  $U_2(A); U_2(B);$

Commit

Aka cascading aborts

# A New Problem: Non-recoverable Schedule



T1

$L_1(A); L_1(B);$  READ(A)

$A := A + 100$

WRITE(A);  $U_1(A)$

READ(B)

$B := B + 100$

WRITE(B);  $U_1(B)$

T2

$L_2(A);$  READ(A)

$A := A * 2$

WRITE(A);

$L_2(B);$  BLOCKED...

...GRANTED; READ(B)

$B := B * 2$

WRITE(B);  $U_2(A); U_2(B);$

Commit

Rollback

Elements A, B written  
by T1 are restored  
to their original value.

# A New Problem: Non-recoverable Schedule



T1

$L_1(A); L_1(B); \text{READ}(A)$   
 $A := A + 100$   
 $\text{WRITE}(A); U_1(A)$

$\text{READ}(B)$   
 $B := B + 100$   
 $\text{WRITE}(B); U_1(B);$

Rollback

Elements A, B written  
by T1 are restored  
to their original value.

T2

$L_2(A); \text{READ}(A)$   
 $A := A * 2$   
 $\text{WRITE}(A);$   
 $L_2(B); \text{BLOCKED...}$

...GRANTED;  $\text{READ}(B)$   
 $B := B * 2$   
 $\text{WRITE}(B); U_2(A); U_2(B);$   
 $\text{Commit}$

Dirty reads of  
A, B lead to  
incorrect writes.

# A New Problem: Non-recoverable Schedule



T1

$L_1(A); L_1(B); \text{READ}(A)$   
 $A := A + 100$   
 $\text{WRITE}(A); U_1(A)$

$\text{READ}(B)$   
 $B := B + 100$   
 $\text{WRITE}(B); U_1(B);$

Rollback

Elements A, B written  
by T1 are restored  
to their original value.

T2

$L_2(A); \text{READ}(A)$   
 $A := A * 2$   
 $\text{WRITE}(A);$   
 $L_2(B); \text{BLOCKED...}$

...GRANTED;  $\text{READ}(B)$   
 $B := B * 2$   
 $\text{WRITE}(B); U_2(A); U_2(B);$   
Commit

Can no longer undo!



# Strict 2PL



The Strict 2PL rule:

All locks are held until commit/abort:

All unlocks are done together with commit/abort.

With strict 2PL, we will get schedules that  
are both conflict-serializable and recoverable

# Strict 2PL



T1

L<sub>1</sub>(A); READ(A)

A := A+100

WRITE(A);

L<sub>1</sub>(B); READ(B)

B := B+100

WRITE(B);

Rollback & U<sub>1</sub>(A); U<sub>1</sub>(B);

T2

L<sub>2</sub>(A); BLOCKED...

...GRANTED; READ(A)

A := A\*2

WRITE(A);

L<sub>2</sub>(B); READ(B)

B := B\*2

WRITE(B);

Commit & U<sub>2</sub>(A); U<sub>2</sub>(B);

# Strict 2PL



- Lock-based systems always use strict 2PL
- Easy to implement: *Implementation*
  - Before a transaction reads or writes an element A, insert an L(A)
  - When the transaction commits/aborts, then release all locks
- Ensures both conflict serializability and recoverability

## Another problem: Deadlocks



- $T_1$ : R(A), W(B)
  - $T_2$ : R(B), W(A)
- 
- $T_1$  holds the lock on A, waits for B
  - $T_2$  holds the lock on B, waits for A

$T_1$  waits for  $T_2$  while  
 $T_2$  waits for  $T_1$

This is a **deadlock!**



# Deadlock Prevention

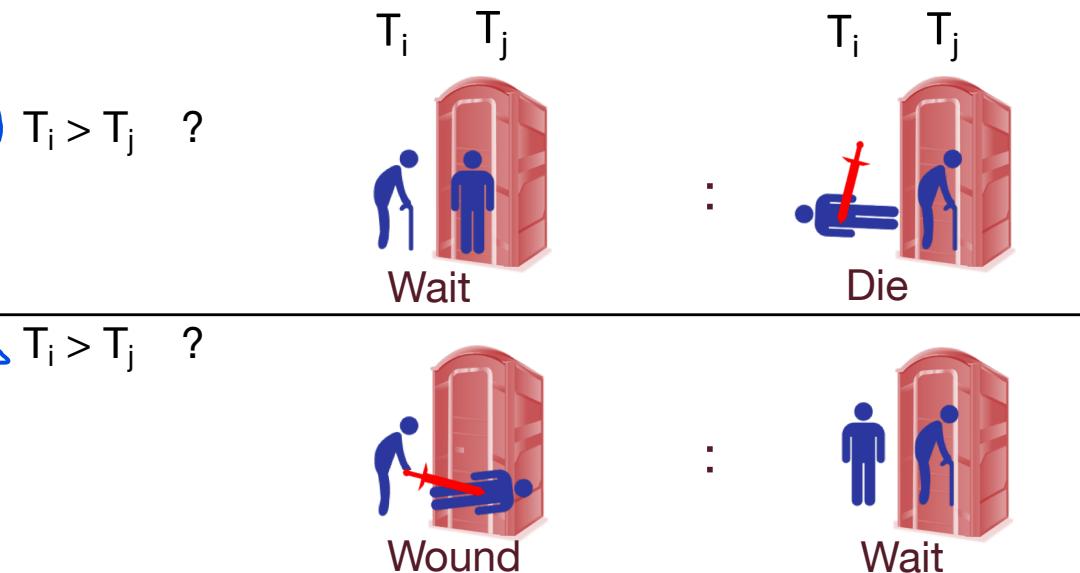


- **Common technique in operating systems**
- **Standard approach: resource ordering**
  - Screen < Network Card < Printer
- **Why is this problematic for transactions?**
  - What order would you impose?

# Deadlock Avoidance



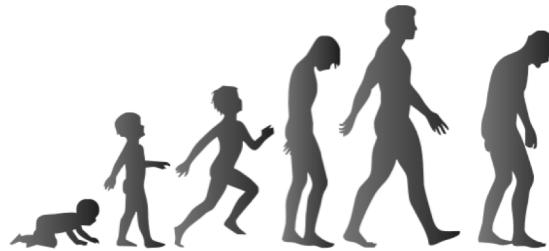
- Assign priorities based on age: (now – start\_time).
- Say  $T_i$  wants a lock that  $T_j$  holds. Two possible policies:
  - Wait-Die: If  $T_i$  has higher priority,  $T_i$  waits for  $T_j$ ; else  $T_i$  aborts
  - Wound-Wait: If  $T_i$  has higher priority,  $T_j$  aborts; else  $T_i$  waits
- Read each of these like a ternary operator (C/C++/java/javascript)



# Deadlock Avoidance: Analysis



- Q: Why do these schemes guarantee no deadlocks?
    - Q: What do the previous images have in common?
  - Important Detail: If a transaction re-starts, make sure it gets its original timestamp. Why?
  - Note: other priority schemes make sense
    - E.g. measures of resource consumption, like #locks acquired
- else might always be  
low priority*



# Deadlock Detection



- Create and maintain a **“waits-for” graph**
- Periodically check for cycles in a graph

# Deadlock Detection, Part 2

**Example:**

**T1:**

**T2:**

**T3:**

**T4:**



T1

T2

T4

T3

# Deadlock Detection, Part 3

**Example:**

$L_1(A)$

**T1:**  $R(A)$

**T2:**

**T3:**

**T4:**



T1

T2

T4

T3

# Deadlock Detection, Part 4

**Example:**

**T1:** R(A) R(D)

**T2:**

**T3:**

**T4:**



T1

T2

T4

T3

# Deadlock Detection, Part 5

**Example:**

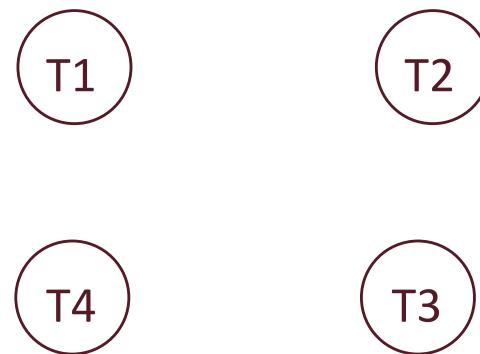
$L_1 \subset D$

T1: R(A) R(D)  $L_2 \subset B$

T2: W(B)

T3:

T4:



# Deadlock Detection, Part 6

**Example:**

wait for  $V_2(LB)$   
↓

T1: R(A) R(D) R(B)

T2: W(B)

T3:

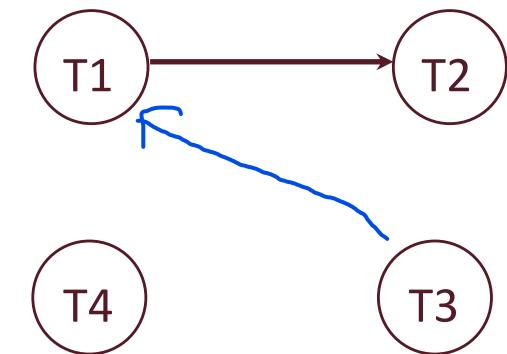
T4:



# Deadlock Detection, Part 7

**Example:**

T1: R(A) R(D)      R(B)      *wait for U, <D)*  
T2:                    W(B)  
T3:  
T4:



# Deadlock Detection, Part 8

**Example:**

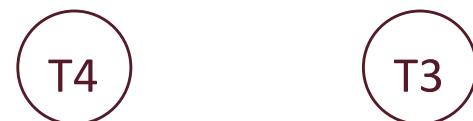
T1: R(A) R(D)

T2: W(B)

T3:

T4:

L(C)



# Deadlock Detection, Part 9

**Example:**

T1: R(A) R(D) R(B)

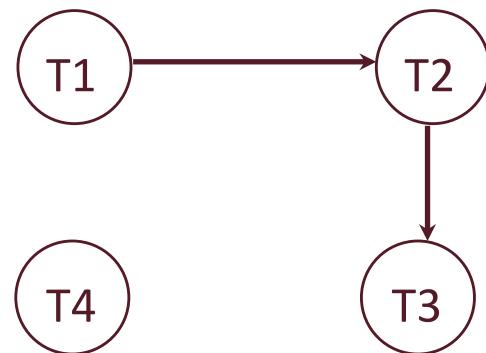
T2: W(B)

T3:

T4:

wait for U<sub>3</sub> <<

W(C)  
R(D) R(C)



# Deadlock Detection, Part 10

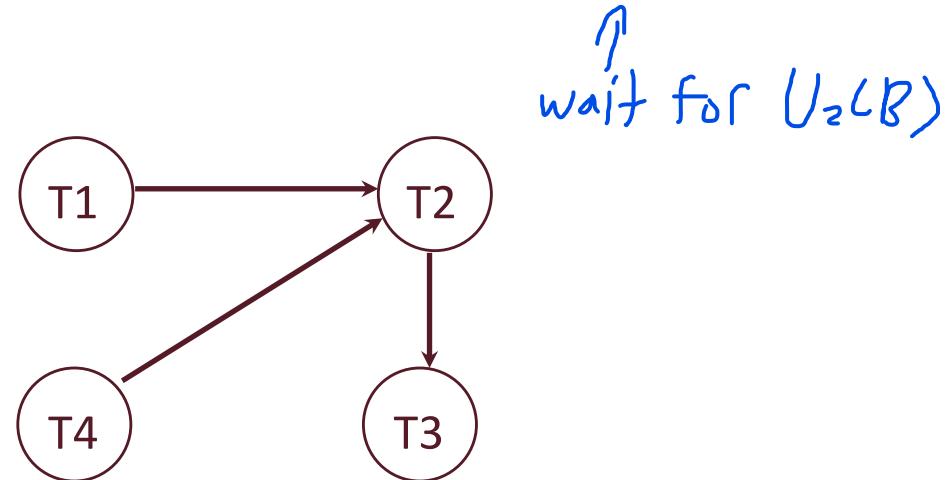
**Example:**

T1: R(A) R(D) R(B)

T2: W(B) W(C)

T3: R(D) R(C)

T4: W(B)



# Deadlock Detection, Part 11

**Example:**

T1: R(A) R(D) R(B)

T2: W(B)

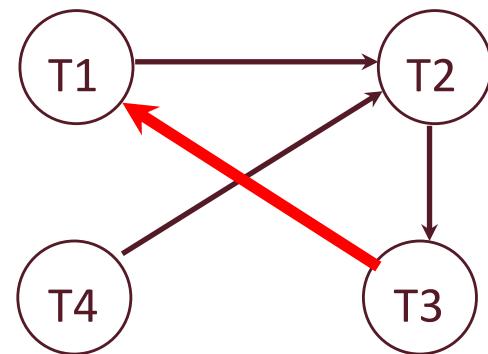
T3: R(D) R(C)

T4:

W(C)

W(B)

wait for U(A)  
W(A) ↙



# Deadlock!



- T1, T2, T3 are deadlocked
  - Doing no good, and holding locks
- T4 still cruising
- In the background, run a deadlock detection algorithm
  - Periodically extract the waits-for graph
  - Find cycles
  - “Shoot” a transaction on the cycle (*kill*)
- Empirical fact
  - Most deadlock cycles are small (2-3 transactions)

# Lock Modes



- **S** = shared lock (for READ)
- **X** = exclusive lock (for WRITE)
- Cannot get new locks after releasing any locks (strict 2PL)

Lock compatibility matrix:

	None	S	X
None			
S			
X			

# Lock Modes



- **S** = shared lock (for READ)
- **X** = exclusive lock (for WRITE)
- Cannot get new locks after releasing any locks (strict 2PL)

Lock compatibility matrix:

	None	S	X
None	✓	✓	✓
S	✓	✓	✗
X	✓	✗	✗

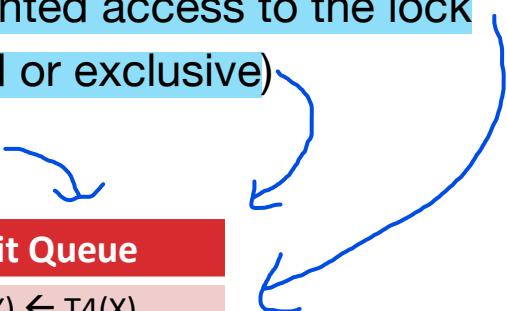
Compatibility matrix

# Lock Management



- Lock and unlock requests handled by Lock Manager
- LM maintains a hashtable, keyed on names of objects being locked.
- LM keeps an entry for each currently held lock
- Entry contains
  - Granted set: Set of txns currently granted access to the lock
  - Lock mode: Type of lock held (shared or exclusive)
  - Wait Queue: Queue of lock requests

	Granted Set	Mode	Wait Queue
A	{T1, T2}	S	T3(X) ← T4(X)
B	{T6}	X	T5(X) ← T7(S)



# Lock Management (continued)



Using compatibility matrix

- **When lock request arrives:**

- Does any txn in Granted Set or Wait Queue want a conflicting lock?
  - If no, put the requester into “granted set” and let them proceed
  - If yes, put requester into wait queue (typically FIFO)

- **Lock upgrade:**

- Txn with shared lock can request to upgrade to exclusive

“Does there exist a conflict between request and a lock request in granted set + wait queue”?

	Granted Set	Mode	Wait Queue
A	{T1, T2}	S	T3(X) ← T4(X)
B	{T6}	X	T5(X) ← T7(S)

# Lock Granularity

what are we putting lock on?



- Fine granularity locking (e.g., tuples)
  - High concurrency
  - High overhead in managing locks
  - E.g., SQL Server
- Coarse grain locking (e.g., tables, entire database)
  - Many false conflicts
  - Less overhead in managing locks
  - E.g., SQL Lite
- Solution: lock escalation changes granularity as needed

locking  
L ( — )

locking

# Lock Granularity, cont



- Hard to decide what granularity to lock
  - Tuples vs Pages vs Tables?
- What is the tradeoff?
  - Fine-grained availability of resources would be nice (e.g. lock per tuple)
  - Small # of locks to manage would also be nice (e.g. lock per table)
  - Can't have both!
    - Or can we???

# Multiple Locking Granularity

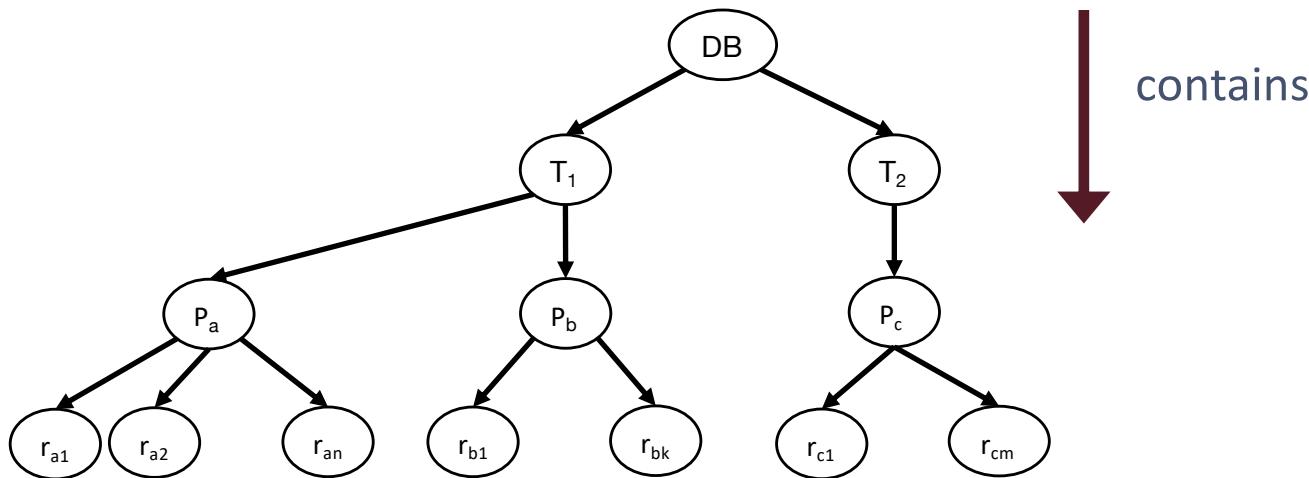


- **Shouldn't have to make same decision for all transactions!**
- Allow data items to be of various sizes
- Define a hierarchy of data granularities, small nested within large
  - Can be represented graphically as a tree.

# Example of Granularity Hierarchy



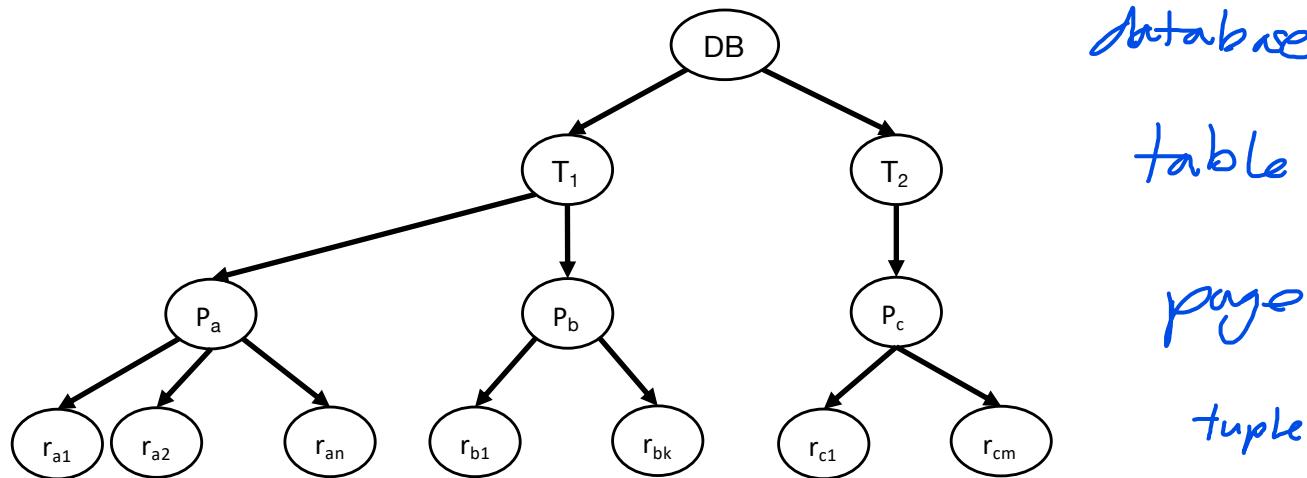
- Data “containers” can be viewed as nested.
- The levels, starting from the coarsest (top) level are
  - Database, Tables, Pages, Records
- When a transaction locks a node in the tree **explicitly**, it **implicitly** locks all the node’s descendants in the same mode.



# Multiple Locking Granularity



- Granularity of locking (level in tree where locking is done):
  - **Fine granularity** (lower in tree): High concurrency, lots of locks (overhead)
  - **Coarse granularity** (higher in tree): Few locks (low overhead), lost concurrency
    - Lost potential concurrency if you don't need everything inside the coarse grain



# Real-World Locking Granularities



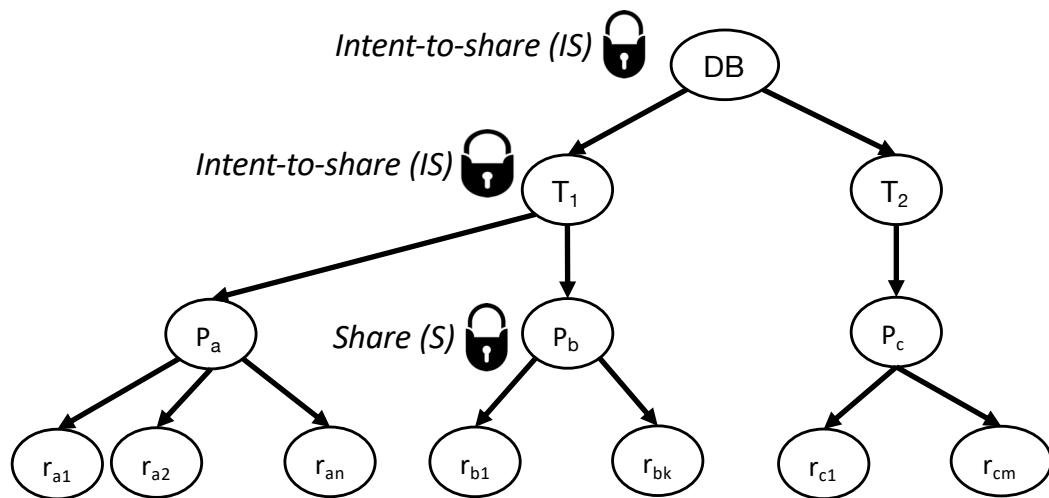
Resource	Description
RID	A row identifier used to lock a single row within a heap.
KEY	A row lock within an index used to protect key ranges in serializable transactions.
PAGE	An 8-kilobyte (KB) page in a database, such as data or index pages.
EXTENT	A contiguous group of eight pages, such as data or index pages.
HoBT	A heap or B-tree. A lock protecting a B-tree (index) or the heap data pages in a table that does not have a clustered index.
TABLE	The entire table, including all data and indexes.
FILE	A database file.
APPLICATION	An application-specified resource.
METADATA	Metadata locks.
ALLOCATION_UNIT	An allocation unit.
DATABASE	The entire database.

From MS SQL Server  
[https://technet.microsoft.com/en-us/library/jj856598\(v=sql.110\).aspx](https://technet.microsoft.com/en-us/library/jj856598(v=sql.110).aspx)

# New Lock Modes and Protocol



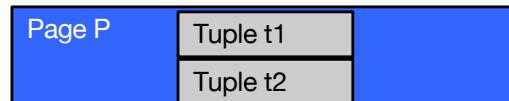
- Allow txns to lock at each level, but with a special protocol using new “intent” locks:
- Before getting S or X lock, txn must have proper intent locks on all its ancestors in the granularity hierarchy.



# New Lock Modes – Intention Lock Modes



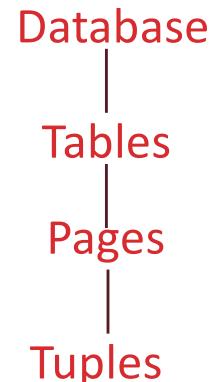
- 3 additional lock modes:
  - **IS:** Intent to get S lock(s) at finer granularity.
  - **IX:** Intent to get X lock(s) at finer granularity.
  - **SIX:** Like S & IX at the same time. Why useful?
- Intention locks allow a higher level node to be locked in S or X mode without having to check all descendant nodes



# Multiple Granularity Locking Protocol



- Each txn starts from the root of the hierarchy.
  - To get S or IS lock on a node, must hold IS or IX on parent node.
  - To get X or IX or SIX on a node, must hold IX or SIX on parent node.
  - Must release locks in bottom-up order.
- 
- Enforce (strict) 2-phase locking as before
  - Protocol is correct in that it is *equivalent to directly setting locks at leaf levels of the hierarchy*.
- 
- What does the lock compatibility matrix look like?



# Lock Compatibility Matrix

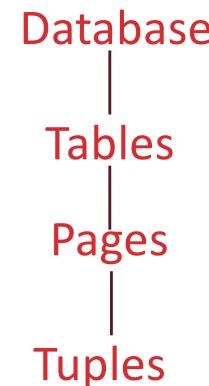
- IS – Intent to get S lock(s) at finer granularity.
- IX – Intent to get X lock(s) at finer granularity.
- SIX mode: Like S & IX at the same time.



Handy simple case to remember:  
Could 2 intent locks be compatible?

Page P	Tuple t1	S	IS
	Tuple t2	X	IX

	IS	IX	S	SIX	X
IS					
IX					
S			true		false
SIX					
X			false		false



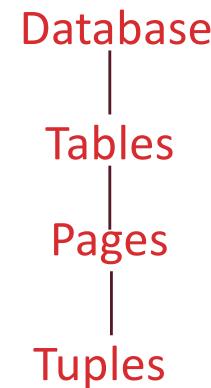
# Lock Compatibility Matrix, Cont

- IS – Intent to get S lock(s) at finer granularity.
- IX – Intent to get X lock(s) at finer granularity.
- SIX mode: Like S & IX at the same time.



Handy simple case to remember:  
Could 2 intent locks be compatible?

	Page P	Tuple t1	S	IS
		Tuple t2	X	IX
IS	true	true	true	false
IX	true	true	false	false
S	true	false	true	false
SIX	true	false	false	false
X	false	false	false	false



	IS	IX	S	SIX	X
IS	true	true	true	true	false
IX	true	true	false	false	false
S	true	false	true	false	false
SIX	true	false	false	false	false
X	false	false	false	false	false

# Real-World Lock Compatibility Matrix



	NL	SCH-S	SCH-M	S	U	X	IS	IU	IX	SIU	SIX	UIX	BU	RS-S	RS-U	RI-N	RI-S	RI-U	RI-X	RX-S	RX-U	RX-X
NL	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N
SCH-S	N	N	C	N	N	N	N	N	N	N	N	I	I	I	I	I	I	I	I	I	I	I
SCH-M	N	C	C	C	C	C	C	C	C	C	C	I	I	I	I	I	I	I	I	I	I	I
S	N	N	C	N	N	C	N	N	C	N	C	C	C	N	N	N	N	C	N	N	C	C
U	N	N	C	N	C	C	N	C	C	C	C	C	C	N	C	N	N	C	C	N	C	C
X	N	N	C	C	C	C	C	C	C	C	C	C	C	C	C	N	C	C	C	C	C	C
IS	N	N	C	N	N	C	N	N	N	N	N	C	I	I	I	I	I	I	I	I	I	I
IU	N	N	C	N	C	C	N	N	N	N	N	C	I	I	I	I	I	I	I	I	I	I
IX	N	N	C	C	C	C	N	N	N	C	C	I	I	I	I	I	I	I	I	I	I	I
SIU	N	N	C	N	C	C	N	N	C	N	C	C	I	I	I	I	I	I	I	I	I	I
SIX	N	N	C	C	C	C	N	N	C	C	C	C	I	I	I	I	I	I	I	I	I	I
UIX	N	N	C	C	C	C	N	C	C	C	C	C	I	I	I	I	I	I	I	I	I	I
BU	N	N	C	C	C	C	C	C	C	C	C	C	N	I	I	I	I	I	I	I	I	I
RS-S	N	I	I	N	N	C	I	I	I	I	I	I	I	N	N	C	C	C	C	C	C	C
RS-U	N	I	I	N	C	C	I	I	I	I	I	I	I	N	C	C	C	C	C	C	C	C
RI-N	N	I	I	N	N	N	I	I	I	I	I	I	I	C	C	N	N	N	N	C	C	C
RI-S	N	I	I	N	N	C	I	I	I	I	I	I	I	C	C	N	N	N	C	C	C	C
RI-U	N	I	I	N	C	C	I	I	I	I	I	I	I	C	C	N	N	C	C	C	C	C
RI-X	N	I	C	C	C	I	I	I	I	I	I	I	C	C	N	C	C	C	C	C	C	C
RX-S	N	I	I	N	N	C	I	I	I	I	I	I	I	C	C	C	C	C	C	C	C	C
RX-U	N	I	I	N	C	C	I	I	I	I	I	I	I	C	C	C	C	C	C	C	C	C
RX-X	N	I	I	C	C	C	I	I	I	I	I	I	C	C	C	C	C	C	C	C	C	C

## Key

N	No Conflict	SIU	Share with Intent Update
I	Illegal	SIX	Shared with Intent Exclusive
C	Conflict	UIX	Update with Intent Exclusive
NL	No Lock	BU	Bulk Update
SCH-S	Schema Stability Locks	RS-S	Shared Range-Shared
SCH-M	Schema Modification Locks	RS-U	Shared Range-Update
S	Shared	RI-N	Insert Range-Null
U	Update	RI-S	Insert Range-Shared
X	Exclusive	RI-U	Insert Range-Update
IS	Intent Shared	RI-X	Insert Range-Exclusive
IU	Intent Update	RX-S	Exclusive Range-Shared
IX	Intent Exclusive	RX-U	Exclusive Range-Update
		RX-X	Exclusive Range-Exclusive

From MS SQL Server

[https://technet.microsoft.com/en-us/library/jj856598\(v=sql.110\).aspx](https://technet.microsoft.com/en-us/library/jj856598(v=sql.110).aspx)

# Phantom Problem



- So far we have assumed the database to be a **static** collection of elements (=tuples)
- If tuples are inserted/deleted then the *phantom problem* appears

Suppose there are two blue products, A1, A2:

# Phantom Problem



T1

```
SELECT *  
FROM Product  
WHERE color='blue'
```

T2

```
INSERT INTO Product(name, color)  
VALUES ('A3','blue')
```

```
SELECT *  
FROM Product  
WHERE color='blue'
```

Is this schedule serializable ?

Suppose there are two blue products, A1, A2:

# Phantom Problem



T1

T2

---

```
SELECT *
FROM Product
WHERE color='blue'
```

```
INSERT INTO Product(name, color)
VALUES ('A3','blue')
```

```
SELECT *
FROM Product
WHERE color='blue'
```

Is this schedule serializable ?

No: T1 sees a “phantom” product A3

Suppose there are two blue products, A1, A2:

# Phantom Problem



T1

T2

---

```
SELECT *
FROM Product
WHERE color='blue'
```

```
INSERT INTO Product(name, color)
VALUES ('A3','blue')
```

```
SELECT *
FROM Product
WHERE color='blue'
```

But this is conflict-serializable!

R<sub>1</sub>(A1);R<sub>1</sub>(A2);W<sub>2</sub>(A3);R<sub>1</sub>(A1);R<sub>1</sub>(A2);R<sub>1</sub>(A3)

W<sub>2</sub>(A3);R<sub>1</sub>(A1);R<sub>1</sub>(A2);R<sub>1</sub>(A1);R<sub>1</sub>(A2);R<sub>1</sub>(A3)

# Phantom Problem



- A “phantom” is a tuple that is invisible during **part** of a transaction execution but not invisible during the **entire** execution
- In our example:
  - T1: reads list of products
  - T2: inserts a new product
  - T1: re-reads: a new product appears !
- Conflict-serializability assumes DB is **static**
- When DB is **dynamic** then c-s is not serializable.



Conflict serializability

# Dealing With Phantoms



- Lock the entire table
- Lock the index entry for ‘blue’
  - If index is available
- Or use predicate locks
  - A lock on an arbitrary predicate

Dealing with phantoms is expensive!

# Summary of Serializability



- Serializable schedule = equivalent to a serial schedule
- (strict) 2PL guarantees *conflict serializability*
  - What is the difference?
- **Static database:**
  - *Conflict serializability* implies serializability
- **Dynamic database:**
  - This no longer holds

# Summary, cont.



- **Correctness criterion for isolation is “serializability”.**
  - In practice, we use “conflict serializability” which is conservative but easy to enforce
- **Two Phase Locking and Strict 2PL: Locks implement the notions of conflict directly**
  - The lock manager keeps track of the locks issued.
  - **Deadlocks** may arise; can either be prevented or detected.
- **Multi-Granularity Locking:**
  - Allows flexible tradeoff between lock “scope” in DB, and # of lock entries in lock table
- **More to the story**
  - Optimistic/Multi-version/Timestamp CC
  - Index “latching”, phantoms
  - Actually, there’s much much more ☺